

Apache Kafka:

Tips and Tricks I wish I had
known before.

Arnaud Esteve
arnaud.esteve@gmail.com





twitter.com/arnaudesteve



github.com/aesteve



linkedin.com/in/arnaud-esteve-74283745

Context

*Past exp. w/ Kafka at **Decathlon***

- + **Confluent** customers' Support + FAQs
- = Lots of stories merged into a “**we**”

Legend



Don't



Be careful



Tip



Rule (of thumb)



Keep in mind



Why would I?

Heard about Kafka?

Ever tried Kafka?

Deployed in production?

server.properties

```
...
auto.topics.create.enable=?
...
...
```

Can you spot the problem?

```
producer = KafkaProducer (value_serializer=lambda m: json.dumps(m).encode('ascii'))  
for customer in customers:  
    producer.send("customer-{}".format(customer),  
                 {"amount": 1234, "date": "2023-01-05T22:02:00.000Z"},  
                 b'orders'  
                 )
```

What did you expect?

The screenshot shows the 'Topics' page from the Confluent Cloud interface. On the left is a sidebar with icons for navigation and management. The main area has a dark header with a search bar, a 'Hide internal topics' dropdown, and a search button. Below is a table with columns for Topic Name, Count, Size, Last Record, Partitions, Factor, and In Sync.

Topics						
Name	Count	Size	Last Record	Partitions	Factor	In Sync
customer-0	≈ 1	127 B	44 seconds ago	1	1	1
customer-1	≈ 1	127 B	44 seconds ago	1	1	1
customer-10	≈ 1	127 B	41 seconds ago	1	1	1
customer-11	≈ 1	127 B	40 seconds ago	1	1	1
customer-12	≈ 1	127 B	40 seconds ago	1	1	1
customer-13	≈ 1	127 B	39 seconds ago	1	1	1
customer-14	≈ 1	127 B	39 seconds ago	1	1	1
customer-15	≈ 1	127 B	39 seconds ago	1	1	1
customer-16	≈ 1	127 B	38 seconds ago	1	1	1
customer-17	≈ 1	127 B	38 seconds ago	1	1	1
customer-18	≈ 1	127 B	37 seconds ago	1	1	1
customer-2	≈ 1	127 B	44 seconds ago	1	1	1
customer-3	≈ 1	127 B	43 seconds ago	1	1	1
customer-4	≈ 1	127 B	43 seconds ago	1	1	1
customer-5	≈ 1	127 B	42 seconds ago	1	1	1
customer-6	≈ 1	127 B	42 seconds ago	1	1	1

What happened

```
for customer in customers:  
    producer = KafkaProducer(value_serializer=lambda m: json.dumps(m).encode('ascii'))  
    producer.send("customer-{}".format(customer),  
                 {"amount": 1234, "date": "2023-01-05T22:02:00.000Z"},  
                 b'orders'  
                 )
```



- ▶ **No** ⇒ `auto.topics.create.enable=false`
(and yes, connectors do create topics...)
- ▶ Human-created? Not necessarily better

💡 **Automate everything** ⇒ integrate in CI/CD

- PRs, reviews, etc. ⇒ Usual IaC workflow
- **Static code checks:**
 - retention policy / period
 - nb. of partitions
 - naming

How should we name our topics?

Can you spot the problem?

```
try (var producer = new KafkaProducer<String, String>(props)) {  
    // legacy topic, before conventions  
    producer.send(new ProducerRecord<>("preprod_customer_orders", "key", "value"));  
    // new topic  
    producer.send(new ProducerRecord<>("preprod.customer.orders", "key", "value"));  
}
```

▶ **Avoid** dynamic parts:
Environment variables,
System properties,
Computed fields (date, etc.)

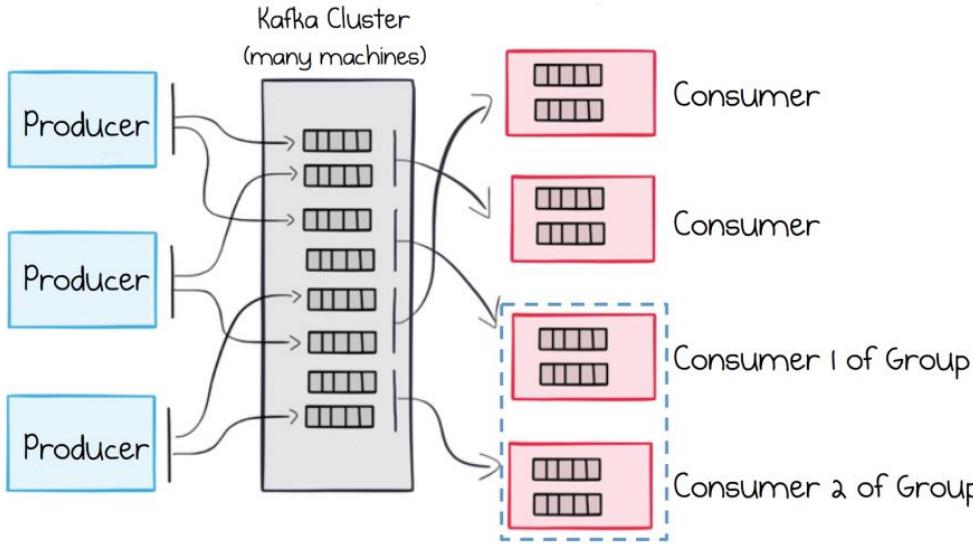
⚠ Dots & Underscores collide!

- 💡 Pick a **convention** & stick to it
 - ⚠ with schema registry conventions
- 💡 Think **Wildcards** for ACLs (preprod.customers.*)
- 💡 **DOCUMENT** everything
 - + automated checks (CI/CD + regex, ...)

👉 <https://medium.com/@kiranprabhu/kafka-topic-naming-conventions-best-practices-6b6b332769a3>

👉 <https://devshawn.com/blog/apache-kafka-topic-naming-conventions/>

What should I use as a key?



Producers spread messages over many partitions, on many machines, where each partition is a little queue. Load balanced consumers (denoted a Consumer Group) share the partitions between them.

- 💡 The producer sets the partition! (=> Partitioner)
- 💡 By default: $\text{hash}(\text{key}) \% \text{partitionCount}$
- Same key \Rightarrow Same partition**



Test.java

```
var topicName = "orders";

try (var admin = AdminClient.create(props)) {
    var topic = new NewTopic(topicName, 15, (short) 1); // 15 partitions
    admin.createTopics(Collections.singleton(topic)).all().get(); // wait for completion
    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());

    try (var producer = new KafkaProducer<String, String>(props)) {
        producer.send(new ProducerRecord<>(topicName, "the-key", "value java"));
    }
}
```

test.py

```
producer = Producer({'bootstrap.servers': 'localhost:9092'})
producer.produce(topic='orders', key='the-key', value='value python')
producer.flush()
```

Wait what?

Topic: orders			
Data	Partitions	Consumer Groups	Configs
Sort: (Oldest) ▾ Partition: (All) ▾ Timestamp: ▾ Search: ▾ Offsets: ▾ Time Format: (RELATIVE) ▾			
Key	Date	Partition	Offset
the-key	31 seconds ago	6	0
value java			
the-key	23 seconds ago	2	0
value python			

librdkafka / CONFIGURATION.md

Top

Preview Code Blame 183 lines (178 loc) · 54.9 KB

Raw    

				<i>type: boolean</i>
partitioner	P	consistent_random	high	Partitioner: random - random distribution, consistent - CRC32 hash of key (Empty and NULL keys are mapped to single partition), consistent_random - CRC32 hash of key (Empty and NULL keys are randomly partitioned), murmur2 - Java Producer compatible Murmur2 hash of key (NULL keys are mapped to single partition), murmur2_random - Java Producer compatible Murmur2 hash of key (NULL keys are randomly partitioned). This is functionally equivalent to the default partitioner in the Java Producer., fnv1a - FNV-1a hash of key (NULL keys are mapped to single partition), fnv1a_random - FNV-1a hash of key (NULL keys are randomly partitioned). <i>Type: string</i>

▶ **AVOID** custom partitioners
(unless you really have to)

▶ **WATCH OUT** for skew:
“One customer makes up for 90% of the traffic”

▶ **AVOID** producing from many languages

⚠ The producer's responsibility
⚠ Keyless can be an option
(if no ordering needed + skew)

- 💡 Partitioning => ordering => causality
- 💡 Determinism for consumers

=> Same key = **causal** to one another

What are cleanup policies?

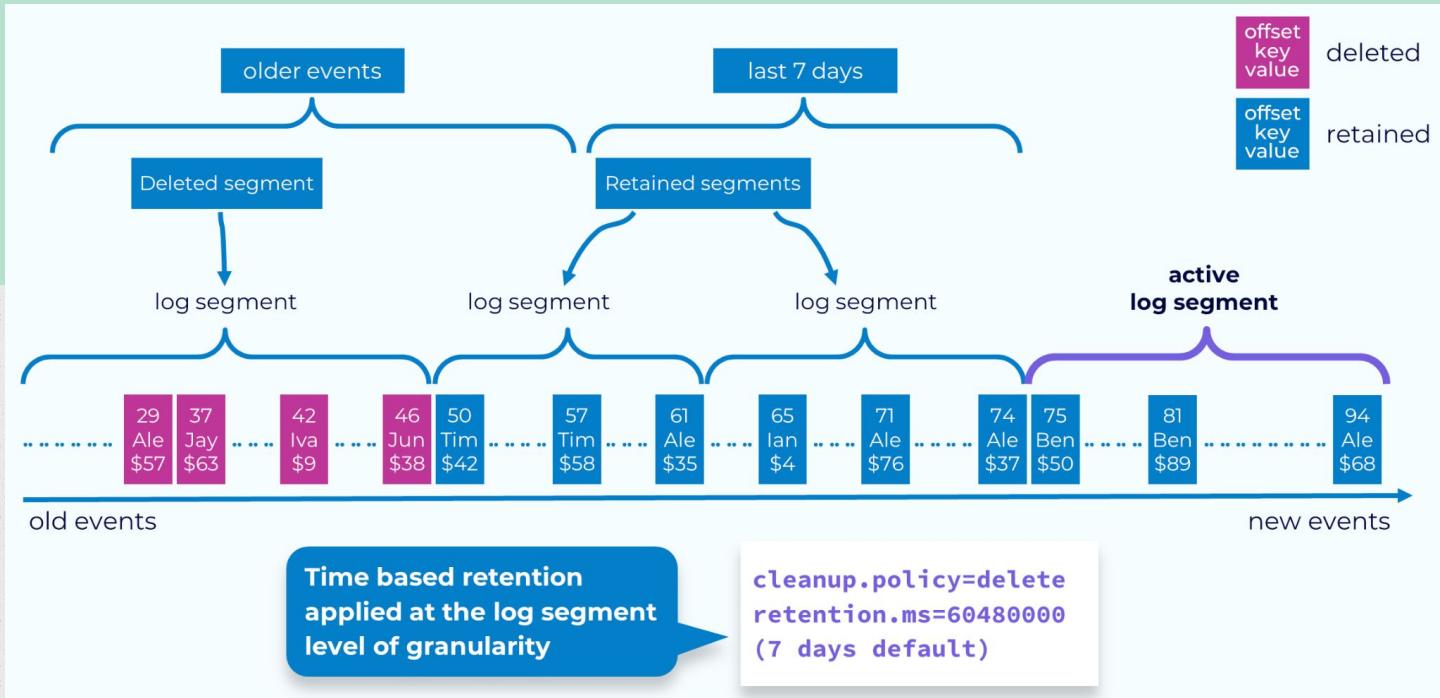
What is going to happen to yesterday's records?

```
try (var admin = AdminClient.create(props)) {
    var topic = new NewTopic("orders", 6, (short) 3);
    topic.configs(Map.of(
        TopicConfig.CLEANUP_POLICY_CONFIG, "compact",
        TopicConfig.RETENTION_MS_CONFIG, "86400000"
    ));
    admin.createTopics(Collections.singleton(topic));
}
```

COMPACT + DELETE

```
try (var admin = AdminClient.create(props)) {
    var topic = new NewTopic("orders", 6, (short) 3);
    topic.configs(Map.of(
        TopicConfig.CLEANUP_POLICY_CONFIG, "compact,delete",
        TopicConfig.RETENTION_MS_CONFIG, "86400000"
    ));
    admin.createTopics(Collections.singleton(topic));
}
```

But still, watch out!



Let's tune the log cleaner, then!

```
try (var admin = AdminClient.create(props)) {
    // 1 partition (easier to observe)
    var topic = new NewTopic(topicName, 1, (short) 1);
    topic.configs(Map.of(
        TopicConfig.CLEANUP_POLICY_CONFIG, "compact,delete",
        TopicConfig.RETENTION_MS_CONFIG, "86400",
        TopicConfig.SEGMENT_MS_CONFIG, "100" // What a good idea 😊
    ));
    admin.createTopics(Collections.singleton(topic)).all().get();
}

props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
try (var producer = new KafkaProducer<String, String>(props)) {
    var i = 0;
    while (++i < 1_000) {
        producer.send(new ProducerRecord<>(topicName, "order-id-" + i, "the order"));
        Thread.sleep(100);
    }
}
```

Yikes...

```
[root@broker appuser]# lsof | wc -l  
22606  
[root@broker appuser]# lsof | wc -l  
31119  
[root@broker appuser]# lsof | wc -l  
32905  
[root@broker appuser]# lsof | wc -l  
41119  
[root@broker appuser]# lsof | wc -l  
43152  
[root@broker appuser]# lsof | wc -l  
46143  
[root@broker appuser]# lsof | wc -l  
47977  
[root@broker appuser]# lsof | wc -l  
56076
```

```
[root@broker appuser]# ls /tmp/kraft-combined-logs/orders-0/  
000000000000000000001998.index  
000000000000000000002031.timeindex  
000000000000000000001998.log  
000000000000000000002032.index  
000000000000000000001998.snapshot  
000000000000000000002032.log  
000000000000000000001998.timeindex  
000000000000000000002032.snapshot  
000000000000000000001999.index  
000000000000000000002032.timeindex  
000000000000000000001999.log  
000000000000000000002033.index  
000000000000000000001999.snapshot  
000000000000000000002033.log  
000000000000000000001999.timeindex  
000000000000000000002033.snapshot  
000000000000000000002000.index  
000000000000000000002033.timeindex  
...
```

▶ **DO NOT**

overtune the log.cleaner

▶ **WATCH OUT** for anti-patterns:

“We must not have any duplicated data”

“We want to reduce the consumption time to ...”

“We want to reduce topic size to ...”

- ⚠ Pick the right `cleanup.policy`
- ⚠ Remember the active segment
- ⚠ Choose the right `retention.period`

💡 `cleanup.policy`:

compact / delete or **BOTH!**

💡 `retention.ms`:

It's in the name!

💡 Compaction is an optimistic mechanism

💡 Use a **StateStore** or a **cache** if need be

How many partitions?

Unit of scalability:

more partitions \Rightarrow more parallelism

Signals to look for:

batch reprocessing "from the origin of time"

very slow / highly parallelizable consumption:
foreach(record) { doSomethingSlow() }

Tradeoffs:

more distributed \Rightarrow more leaders \Rightarrow longer rebalance

Also, **less-efficient batching!**



Highly composite numbers:

6, 12, 24, 48, 60, 120, 180, 240, 360

\rightarrow divide workload more evenly



<https://www.confluent.io/blog/how-choose-number-topics-partitions-kafka-cluster>

How can I know upfront?

And what if I mess up?

No problem!

```
var topicName = "orders";
var key = UUID.randomUUID().toString();
try (var admin = AdminClient.create(props)) {
    // 3 partitions originally
    var topic = new NewTopic(topicName, 3, (short) 1);
    admin.createTopics(Collections.singleton(topic)).all().get(); // wait for completion

    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, LongSerializer.class.getName());
    try (var producer = new KafkaProducer<String, Long>(props)) {
        producer.send(new ProducerRecord<>(topicName, key, 1L));
    }
    admin.createPartitions(Map.of(topicName, NewPartitions.increaseTo(10))).all().get();
    try (var producer = new KafkaProducer<String, Long>(props)) {
        producer.send(new ProducerRecord<>(topicName, key, 2L));
    }
}
```

No problem, really?

Topic: orders

Data Partitions Consumer Groups Configs ACLS Logs

Sort: (Oldest) ▾ Partition: (All) ▾ Timestamp: ▾ Search: ▾ Offsets: ▾ Time Format: (RELATIVE) ▾

Key	Date	Partition	Offset
7e80b30b-a929-458f-9f2b-ae634a637641	33 seconds ago	0	0
7e80b30b-a929-458f-9f2b-ae634a637641	33 seconds ago	6	0

Same key  different partitions 

▶ **DO NOT** increase the nb. of partitions

Unless you know what you're doing:

=> you don't care about ordering

=> no stateful consumers



Create a **new topic** with more partitions instead

=> replicate data to new topic

=> deploy the use-case requiring more partitions

=> migrate consumers progressively

=> migrate producers

CLIENTS

Can you spot the error?

```
var props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("ack", "all"); // I need the strongest delivery guarantee!!
props.put("key.serializer", StringSerializer.class.getName());
props.put("value.serializer", StringSerializer.class.getName());
var topicName = "orders";
try (var producer = new KafkaProducer<String, String>(props)) {
    producer.send(new ProducerRecord<String, String>(topicName, "the-key", "some value"), (metadata, ex) -> {
        LOG.info(
            "Successfully published to all replicas of partition {} (Offset is {}),"
            + metadata.partition(),
            metadata.offset()
        );
    });
}
```

Here it is!

ackS

```
var props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("ack", "all");
props.put("key.serializer", StringSerializer.class.getName());
props.put("value.serializer", StringSerializer.class.getName());
var topicName = "orders";
try (var producer = new KafkaProducer<String, String>(props)) {
    producer.send(new ProducerRecord<String, String>(topicName, "the-key", "some value"), (metadata, ex) -> {
        LOG.info(
            "Successfully published to all replicas of partition {} (Offset is {}),"
            + metadata.partition(),
            metadata.offset()
        );
    });
}
```

Safer this way

```
var props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.ACKS_CONFIG, "all");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG, StringSerializer.class.getName());
var topicName = "orders";
try (var producer = new KafkaProducer<String, String>(props)) {
    producer.send(new ProducerRecord<String, String>(topicName, "the-key", "some value"), (metadata, ex) -> {
        LOG.info(
            "Successfully published to all replicas of partition {} (Offset is {}),"
            + metadata.partition(),
            metadata.offset()
        );
    });
}
```



Constants, constants everywhere

📝 Don't waste time with typos 🙏

💡 Both for keys and values
(EXACTLY_ONCE_V2 , SASL_PLAIN , etc.)

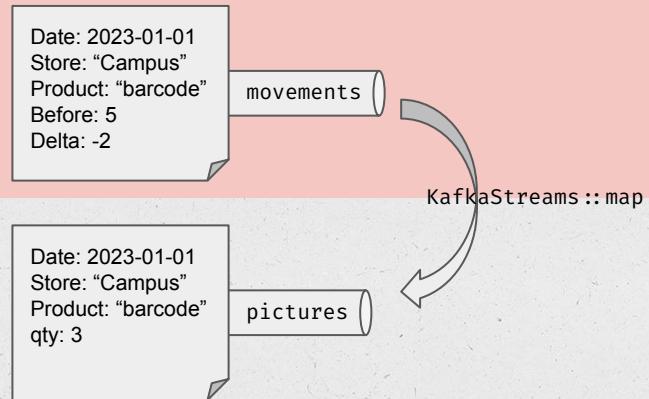
AdminClientConfig
SecurityConfig
ProducerConfig
CommonClientConfig
StreamsConfig
etc.



<https://kafka.apache.org/35/javadoc/constant-values.html>

The Ops strike back

```
var keySerde = Serdes.String();
var movementSerde = new StockMovementSerde();
var pictureSerde = new StockPictureSerde();
var props = new Properties();
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(StreamsConfig.APPLICATION_ID_CONFIG, System.getenv("POD_NAME_VERSION"));
props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
var builder = new StreamsBuilder();
builder.stream("movements", Consumed.with(keySerde, movementSerde))
    .map((key, movement) -> {
        var qtyLeft = movement.before() + movement.delta(); // trivial computation
        return newKeyValue<>(
            key,
            newStockPicture(movement.date(), movement.store(), movement.product(), qtyLeft)
        );
    })
    .to("pictures", Produced.with(keySerde, pictureSerde));
try (var app = new KafkaStreams(builder.build(), props)) {
    app.start();
}
```



Guess what...

```
var keySerde = Serdes.String();
var movementSerde = new StockMovementSerde();
var pictureSerde = new StockPicturesSerde();
var props = new Properties();
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(StreamsConfig.APPLICATION_ID_CONFIG, System.getenv("POD_NAME_VERSION"));
props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
var builder = new StreamsBuilder();
builder.stream("stock-movements", Consumed.with(keySerde, movementSerde))
    .map((key, movement) -> {
        var qtyLeft = movement.before() + movement.delta();
        return newKeyValue<>(
            key,
            newStockPicture(movement.date(), movement.store(), movement.product(), qtyLeft)
        );
    })
    .to("stock-pictures", Produced.with(keySerde, pictureSerde));
try (var app = new KafkaStreams(builder.build(), props)) {
    app.start();
}
```

Changing this

Restarts from here

Mmh this product hasn't
been sold since 2020???

▶ **Avoid** dynamic parts:
Environment variables,
System properties,
Computed fields (date, etc.)

⚠ Consumer Id, Streams Id
⚠ Producers if transactional
Have an id too
⚠ `client.id != group.id`

- 💡 Deterministic identifiers for clients
- 💡 Constant in code, versioned in VCS
- 💡 Potentially extracted at build time

=> the **code drives identifiers**

Wtf why are these records duplicated?

▶ Avoid

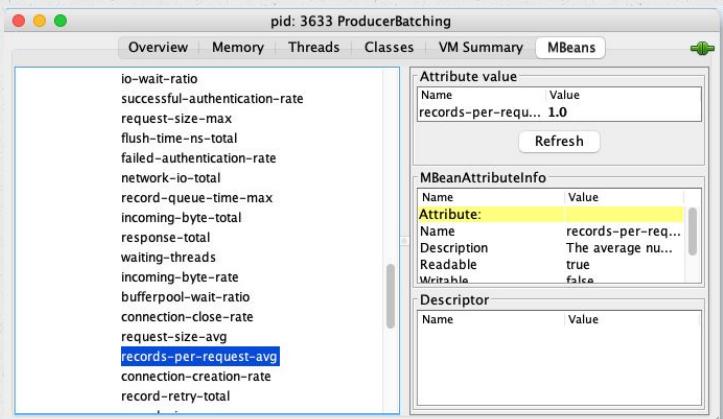
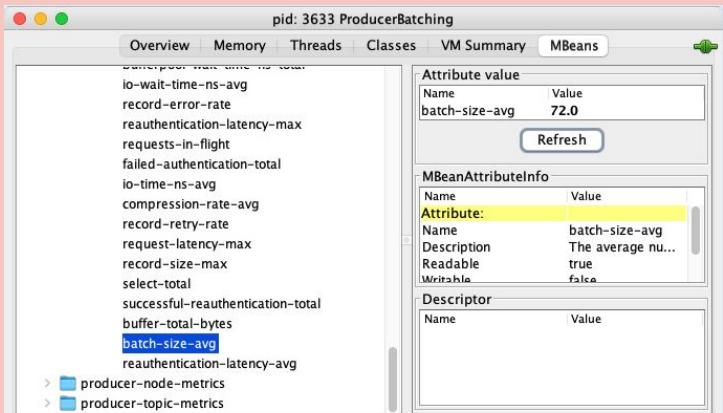
Assuming everything is EO
Setting EOS by default

⚠ Understand delivery semantics
“Know what you’re doing”
⚠ Tradeoff durability / perf.

- 💡 Set **idempotency** on producers
 - 💡 Connectors => check documentation
 - 💡 Kafka Streams, set:
`StreamsConfig.PROCESSING_GUARANTEE_CONFIG`
 - 📏 EOS is no magic, only for Kafka-to-Kafka
- ⇒ design around **At Least Once** where possible

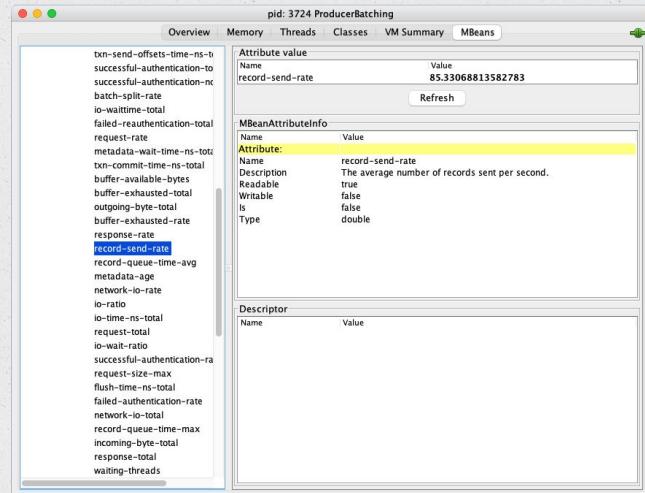
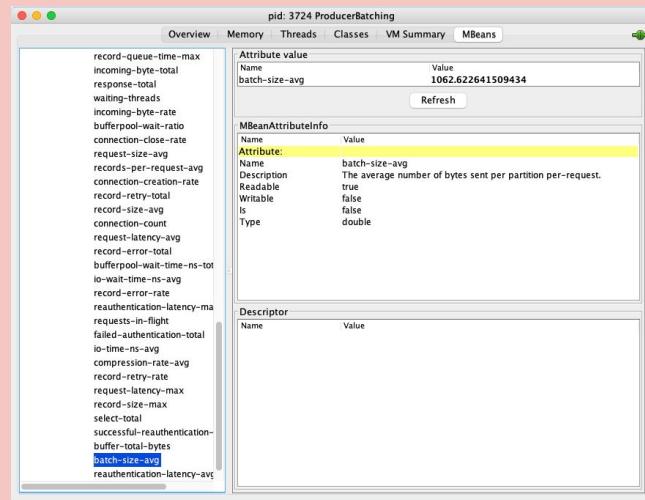
I'm just trying to batch...

```
var props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.BATCH_SIZE_CONFIG, "10000");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
var topic = "some-topic";
try (var producer = new KafkaProducer<String, String>(props)){
    var i = 0;
    while (true) {
        producer.send(new ProducerRecord<>(topic, Integer.toString(++i)));
        Thread.sleep(10);
    }
} catch (InterruptedException e) {
    throw new RuntimeException(e);
}
```



I'm just trying to batch...

```
var props = new Properties ();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.BATCH_SIZE_CONFIG, "10000");
props.put(ProducerConfig.LINGER_MS_CONFIG, "1000");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
var topic = "some-topic";
try (var producer = new KafkaProducer<String, String>(props)) {
    var i = 0;
    while (true) {
        producer.send(new ProducerRecord<>(topic, Integer.toString(++i)));
        Thread.sleep(10);
    }
} catch (InterruptedException e) {
    throw new RuntimeException(e);
}
```



Producer batching + compression

💡 Use JMX Metrics to understand what's going on

💡 Increase your chances of **batching!!!**

`linger.ms > 0`

💡 THEN, configure **compression** (compression acts on batches)

+ saves space on disk

📝 but compress on the producer ⇒ `compression.type = producer`

+ saves bandwidth!

+ saves CPU on brokers ($\sim=$ 0-copy)

Producers: misc



⚠️ Record Timestamp (configurable at topic level)

💡 by default: `CreateTime` (== producer time)

⚠️ can be set to `LogAppendTime` (== broker time)



💡 Use [record headers!!!](#)

OpenTracing

Metadata (like PII, etc.)

Lineage

Which format should I use?



It's on YOU!

Format = YOUR choice



- 尺 What about the rest of your ecosystem? (Protobuf ↔ Google services?)
- 尺 Well-supported in your ecosystem?
 - ⚠ Dynamic languages + Avro?
 - ⚠ Maturity of the library?
- 尺 JSON Schema can be considered the “default goto”
 - ⚠ But isn’t compressed
 - 💡 Avro & Protobuf => generate code from schemas

Avro-specific

- 💡 Specific records, almost always
 - ⚠ Consider generic records as a smell (but...)



Sure... Sure... Really useful...

What matters is the Schema!

Use & Manage Schemas 😊

e.g. by using the Schema Registry

Compatibility Type	Changes allowed	Check against which schemas	Upgrade first
BACKWARD	<ul style="list-style-type: none"> Delete fields Add optional fields 	Last version	Consumers
BACKWARD_TRANSITIVE	<ul style="list-style-type: none"> Delete fields Add optional fields 	All previous versions	Consumers
FORWARD	<ul style="list-style-type: none"> Add fields Delete optional fields 	Last version	Producers
FORWARD_TRANSITIVE	<ul style="list-style-type: none"> Add fields Delete optional fields 	All previous versions	Producers
FULL	<ul style="list-style-type: none"> Add optional fields Delete optional fields 	Last version	Any order
FULL_TRANSITIVE	<ul style="list-style-type: none"> Add optional fields Delete optional fields 	All previous versions	Any order
NONE	<ul style="list-style-type: none"> All changes are accepted 	Compatibility checking disabled	Depends

Compatibility modes:

- 💡 **fORward** ⇒ upgrade **pROducer** first
- 💡 **baCKward** ⇒ upgrade **cONsumers** first

🚩 **Avoid** the non transitive ones
(esp. **If big retention period**)

💡 As conservative as possible
💡 Default setting at SR level, be drastic:
 尺 set to **FULL_TRANSITIVE**
 尺 then override per topic if really needed



Why is Backward the default?

Important

The Confluent Schema Registry default compatibility type is `BACKWARD`, not `BACKWARD_TRANSITIVE`.

The main reason that `BACKWARD` compatibility mode is the default, and preferred for Kafka, is so that you can rewind consumers to the beginning of the topic. With `FORWARD` compatibility mode, you aren't guaranteed the ability to read old messages.

Also `FORWARD` compatibility mode is harder to work with. In a sense, you need to anticipate all future changes. For example, in `FORWARD` compatibility mode with Protobuf, you cannot add new message types to a schema.

Can I test it?



Unit Tests: think Property-Based Testing

`deserializeWithV2(serializeWithV1(record))` ← checks backwards compatibility

`deserializeWithV1(serializeWithV2(record))` ← checks forward compatibility

Inverse property: `deserialize(serialized(record)) = record`



Use: [io.confluent.kafka.schemaregistry.CompatibilityChecker](#) (schema-registry lib)

```
78     // visible for testing ←
79     public List<String> isCompatible(
80         ParsedSchema newSchema, List<? extends ParsedSchema> previousSchemas
81     ) {
82         List<? extends ParsedSchema> previousSchemasCopy = new ArrayList<>(previousSchemas);
83         // Validator checks in list order, but checks should occur in reverse chronological order
84         Collections.reverse(previousSchemasCopy);
85         return validator.validate(newSchema, previousSchemasCopy);
86     }
--
```



Integ. Testing / Schema Promotion



⇒ leverage the Schema Registry [REST API](#), or [maven plugin](#)



Speaking of testing...

Testing

- ▶ **Avoid** integ. testing only
- ▶ **Avoid** using the same topics lacrosse tests

Integ. Testing: `docker forAll(testSuite) (in CI)`:

- ⚠ topic-naming (UUID) strategy per test
- ⚠ **FLUSH** if test involves a real producer!

Simple, unit testing

- 💡 [MockProducer](#) / [MockConsumer](#)
(+ whatever your framework provides)
- 💡 [TopologyTestDriver](#) for streams
↳ really awesome
librdkafka <= rdkafka-mock

Advanced, Integ. Testing

- 💡 [EmbeddedKafkaCluster](#)
- 💡 TestContainers
 - kafka, redpanda, or docker-compose!
(if you need the whole stack)
 - 💡 KRaft (KIP-500) is out!! (zookeeper less)



I need to test performances

Let me tell you a story about A yellow fish with black stripes.

Perf. testing

- ▶ **DO NOT** use custom scripts to create load
- ▶ **DO NOT** test blindly
- ▶ **DO NOT** set random expectations, please

⚠ Performances in Kafka is a tradeoff between:

Latency

Bandwidth

HEAVILY depends on your workload

尺 **MEASURE**, don't guess

Use proper tooling

- 尺 JMeter, Gatling, Trogdor

Testing a cluster?

kafka-producer-perf-test.sh

kafka-consumer-perf-test.sh

💡 TestContainers

kafka, redpanda, or docker-compose!
(if you need the whole stack)

💡 KRaft (KIP-500) is out!! (zookeeper less)

💡 **MIRROR** your traffic as much as you can

💡 eg: MirrorMaker 2

尺 Same:

Traffic **Shape**

Repartition (`topic*partitions`)
Message format

Chaos Testing

- ▶ **AVOID** the rabbit hole
- ▶ **DO NOT** set unrealistic expectations

- 💡 [Trogdor](#) for Kafka workloads with faults
- 💡 Testcontainers + [Toxiproxy](#)
- 💡 Also: Docker + [Pumba](#)

Keep it Simple

We must not lose any message

We must recover in time

Check exception handling!!

ex1: check the consumer lag shape ⇒ does it match the outage?

ex2: just look for unexpected behaviour:

10s broker downtime leads to 99% CPU usage on consumers for 10min

Any architectural advice?

Kafka Connect



Remember it's a library!

- Drives development
 - What do you do in case of failure?
 - Where do you restart from (source)
 - How do you know if the record is written? (sink) When do you ACK?
 - Start asking yourself those questions?
 - Using Connect as a lib?
- **What does it bring?**
 - Deployment model (worker / tasks)
 - Metrics & all
 - Benefit from an existing infra? (connect instance running already?)
 - SerDe agnostic!!!
 - SMTs!!
- **Examples?**
 - Polling from an external service (HTTP, example: OpenSky API)
 - Protocol-Bridging (websockets to Kafka, signalR to Kafka, ...)

Questions?

In the director's cut: (90+ min)

- 💡 **linger.ms** & batching
- 💡 A fun (😅) story about **performance testing** and 🐝
- 💡 Ramblings re. **Idempotency** and Exactly Once Semantics
- 💡 Dealing w/ configuration

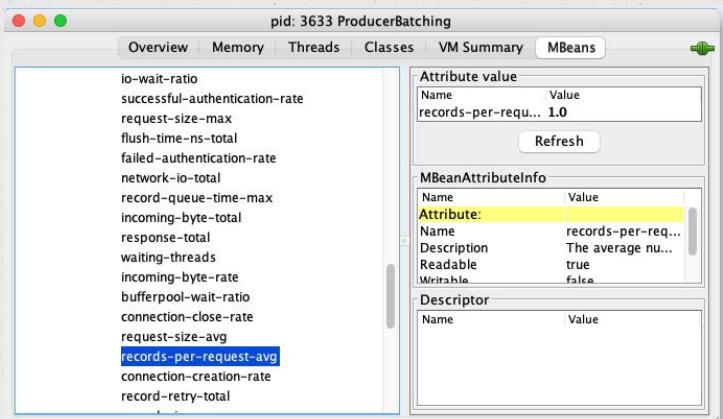
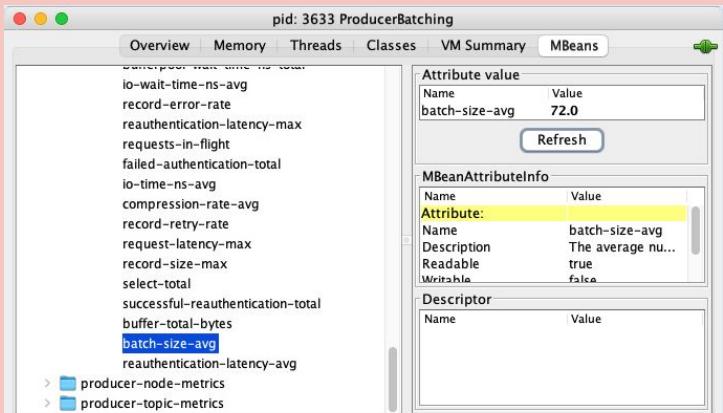


Arnaud Esteve
arnaud.esteve@gmail.com



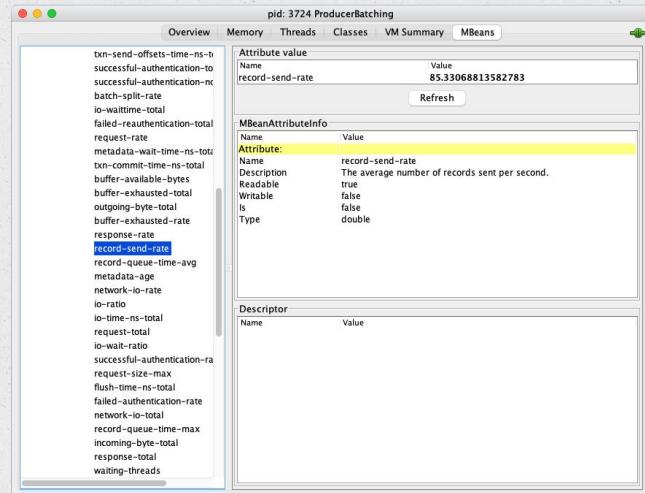
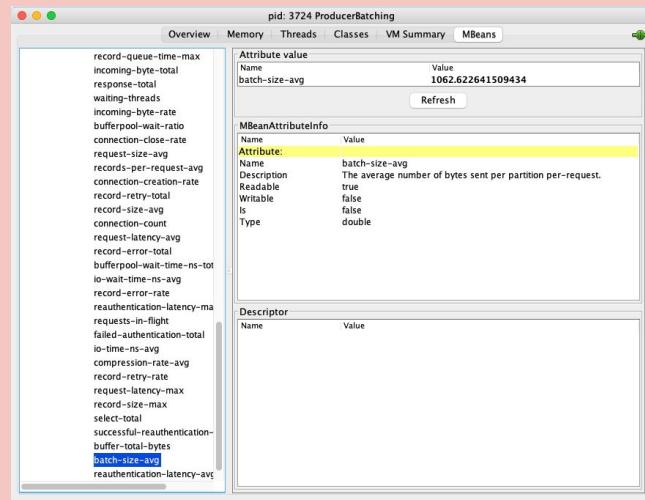
I'm just trying to batch...

```
var props = new Properties();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.BATCH_SIZE_CONFIG, "10000");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName());
var topic = "some-topic";
try (var producer = new KafkaProducer<String, String>(props)){
    var i = 0;
    while (true) {
        producer.send(new ProducerRecord<>(topic, Integer.toString(++i)));
        Thread.sleep(10);
    }
} catch (InterruptedException e) {
    throw new RuntimeException(e);
}
```



I'm just trying to batch...

```
var props = new Properties ();
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.BATCH_SIZE_CONFIG, "10000");
props.put(ProducerConfig.LINGER_MS_CONFIG, "1000");
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName ());
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class.getName ());
var topic = "some-topic";
try (var producer = new KafkaProducer<String, String>(props)) {
    var i = 0;
    while (true) {
        producer.send(new ProducerRecord<>(topic, Integer.toString(++i)));
        Thread.sleep(10);
    }
} catch (InterruptedException e) {
    throw new RuntimeException (e);
}
```



Producer batching + compression

💡 Use JMX Metrics to understand what's going on

💡 Increase your chances of **batching!!!**

`linger.ms > 0`

💡 THEN, configure **compression** (compression acts on batches)

+ saves space on disk

📝 but compress on the producer ⇒ `compression.type = producer`

+ saves bandwidth!

+ saves CPU on brokers ($\sim=$ 0-copy)

Questions?

In the director's cut

- 💡 **linger.ms** & batching
- 💡 Storing configuration
- 💡 A story about **performance testing** and 🐝
- 💡 Ramblings re. Idempotency and Exactly Once Semantics

Arnaud Esteve
arnaud.esteve@gmail.com





I need to test performances

Let me tell you a story about A yellow fish with black stripes.

Perf. testing

- ▶ **DO NOT** use custom scripts to create load
- ▶ **DO NOT** test blindly
- ▶ **DO NOT** set random expectations, please

⚠ Performances in Kafka is a tradeoff between:

Latency

Bandwidth

HEAVILY depends on your workload

尺 **MEASURE**, don't guess

Use proper tooling

- 尺 JMeter, Gatling, Trogdor

Testing a cluster?

kafka-producer-perf-test.sh

kafka-consumer-perf-test.sh

💡 TestContainers

kafka, redpanda, or docker-compose!
(if you need the whole stack)

💡 KRaft (KIP-500) is out!! (zookeeper less)

💡 **MIRROR** your traffic as much as you can

💡 eg: MirrorMaker 2

尺 Same:

Traffic **Shape**

Repartition (`topic*partitions`)
Message format

Chaos Testing

- ▶ **AVOID** the rabbit hole
- ▶ **DO NOT** set unrealistic expectations

- 💡 [Trogdor](#) for Kafka workloads with faults
- 💡 Testcontainers + [Toxiproxy](#)
- 💡 Also: Docker + [Pumba](#)

Keep it Simple

We must not lose any message

We must recover in time

Check exception handling!!

ex1: check the consumer lag shape ⇒ does it match the outage?

ex2: just look for unexpected behaviour:

10s broker downtime leads to 99% CPU usage on consumers for 10min

Questions?

In the director's cut

- 💡 **linger.ms** & batching
- 💡 Storing configuration
- 💡 A story about **performance testing** and 🐝
- 💡 Ramblings re. Idempotency and Exactly Once Semantics

Arnaud Esteve
arnaud.esteve@gmail.com



Wtf why are these records duplicated?

▶ Avoid

Assuming everything is EO
Setting EOS by default

⚠ Understand delivery semantics
“Know what you’re doing”
⚠ Tradeoff durability / perf.

- 💡 Set **idempotency** on producers
 - 💡 Connectors => check documentation
 - 💡 Kafka Streams, set:
`StreamsConfig.PROCESSING_GUARANTEE_CONFIG`
 - 📏 EOS is no magic, only for Kafka-to-Kafka
- ⇒ design around **At Least Once** where possible

Questions?

In the director's cut

- 💡 **linger.ms** & batching
- 💡 Storing configuration
- 💡 A story about **performance testing** and 🐝
- 💡 Ramblings re. Idempotency and Exactly Once Semantics

Arnaud Esteve
arnaud.esteve@gmail.com

