

Tales of a Cloud Migration

Using Apache Kafka

0. Context

Some background

- Personal experience: World-Scale Retailer (Decathlon)
 - Cloud Migration
 - Monolith to micro-services
- Confluent Customers:
 - Retail, Banking, IoT, ...
 - Different contexts, timelines

⇒ Merged into a generic “we”

0. Context

Why?

- “We” could have saved so much time
- A lot of common patterns we’ll explore
- What we missed

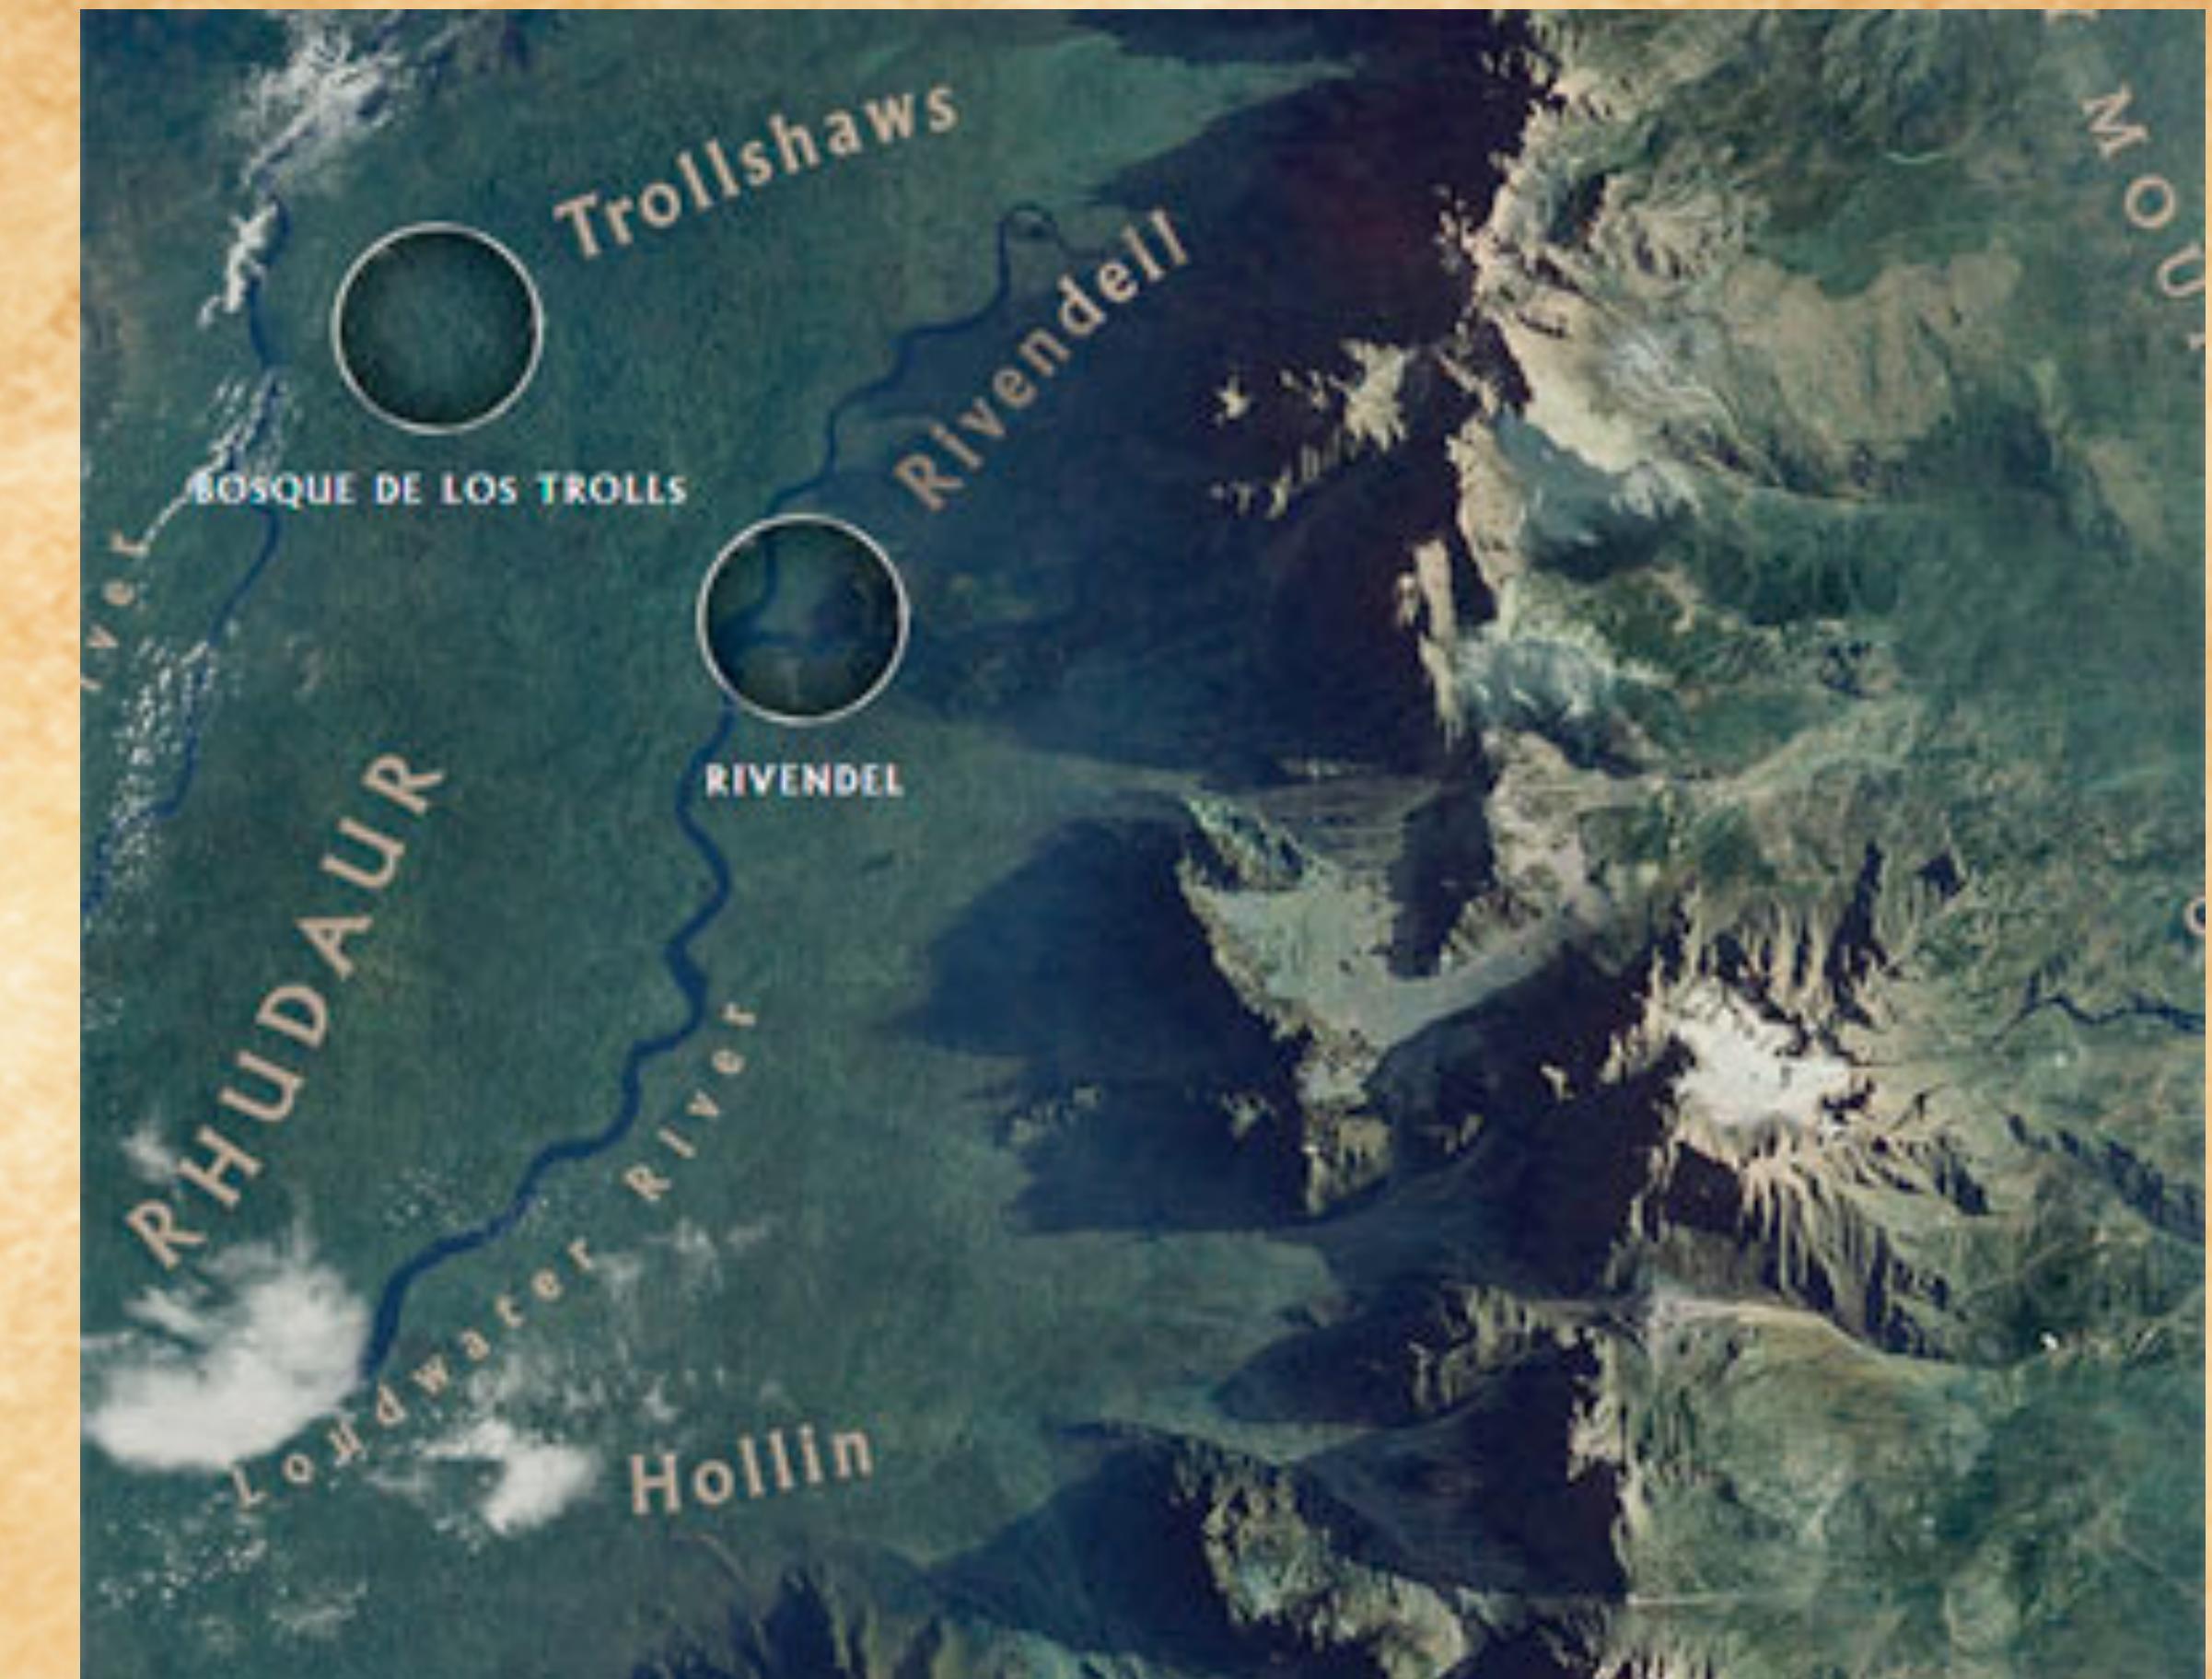
⇒ “What I wish someone told us”

Agenda

1.  Describing the problem(s)
2.  Tackling Confusion
3.  The tactics
4.  The tooling
5.  A word on what was next

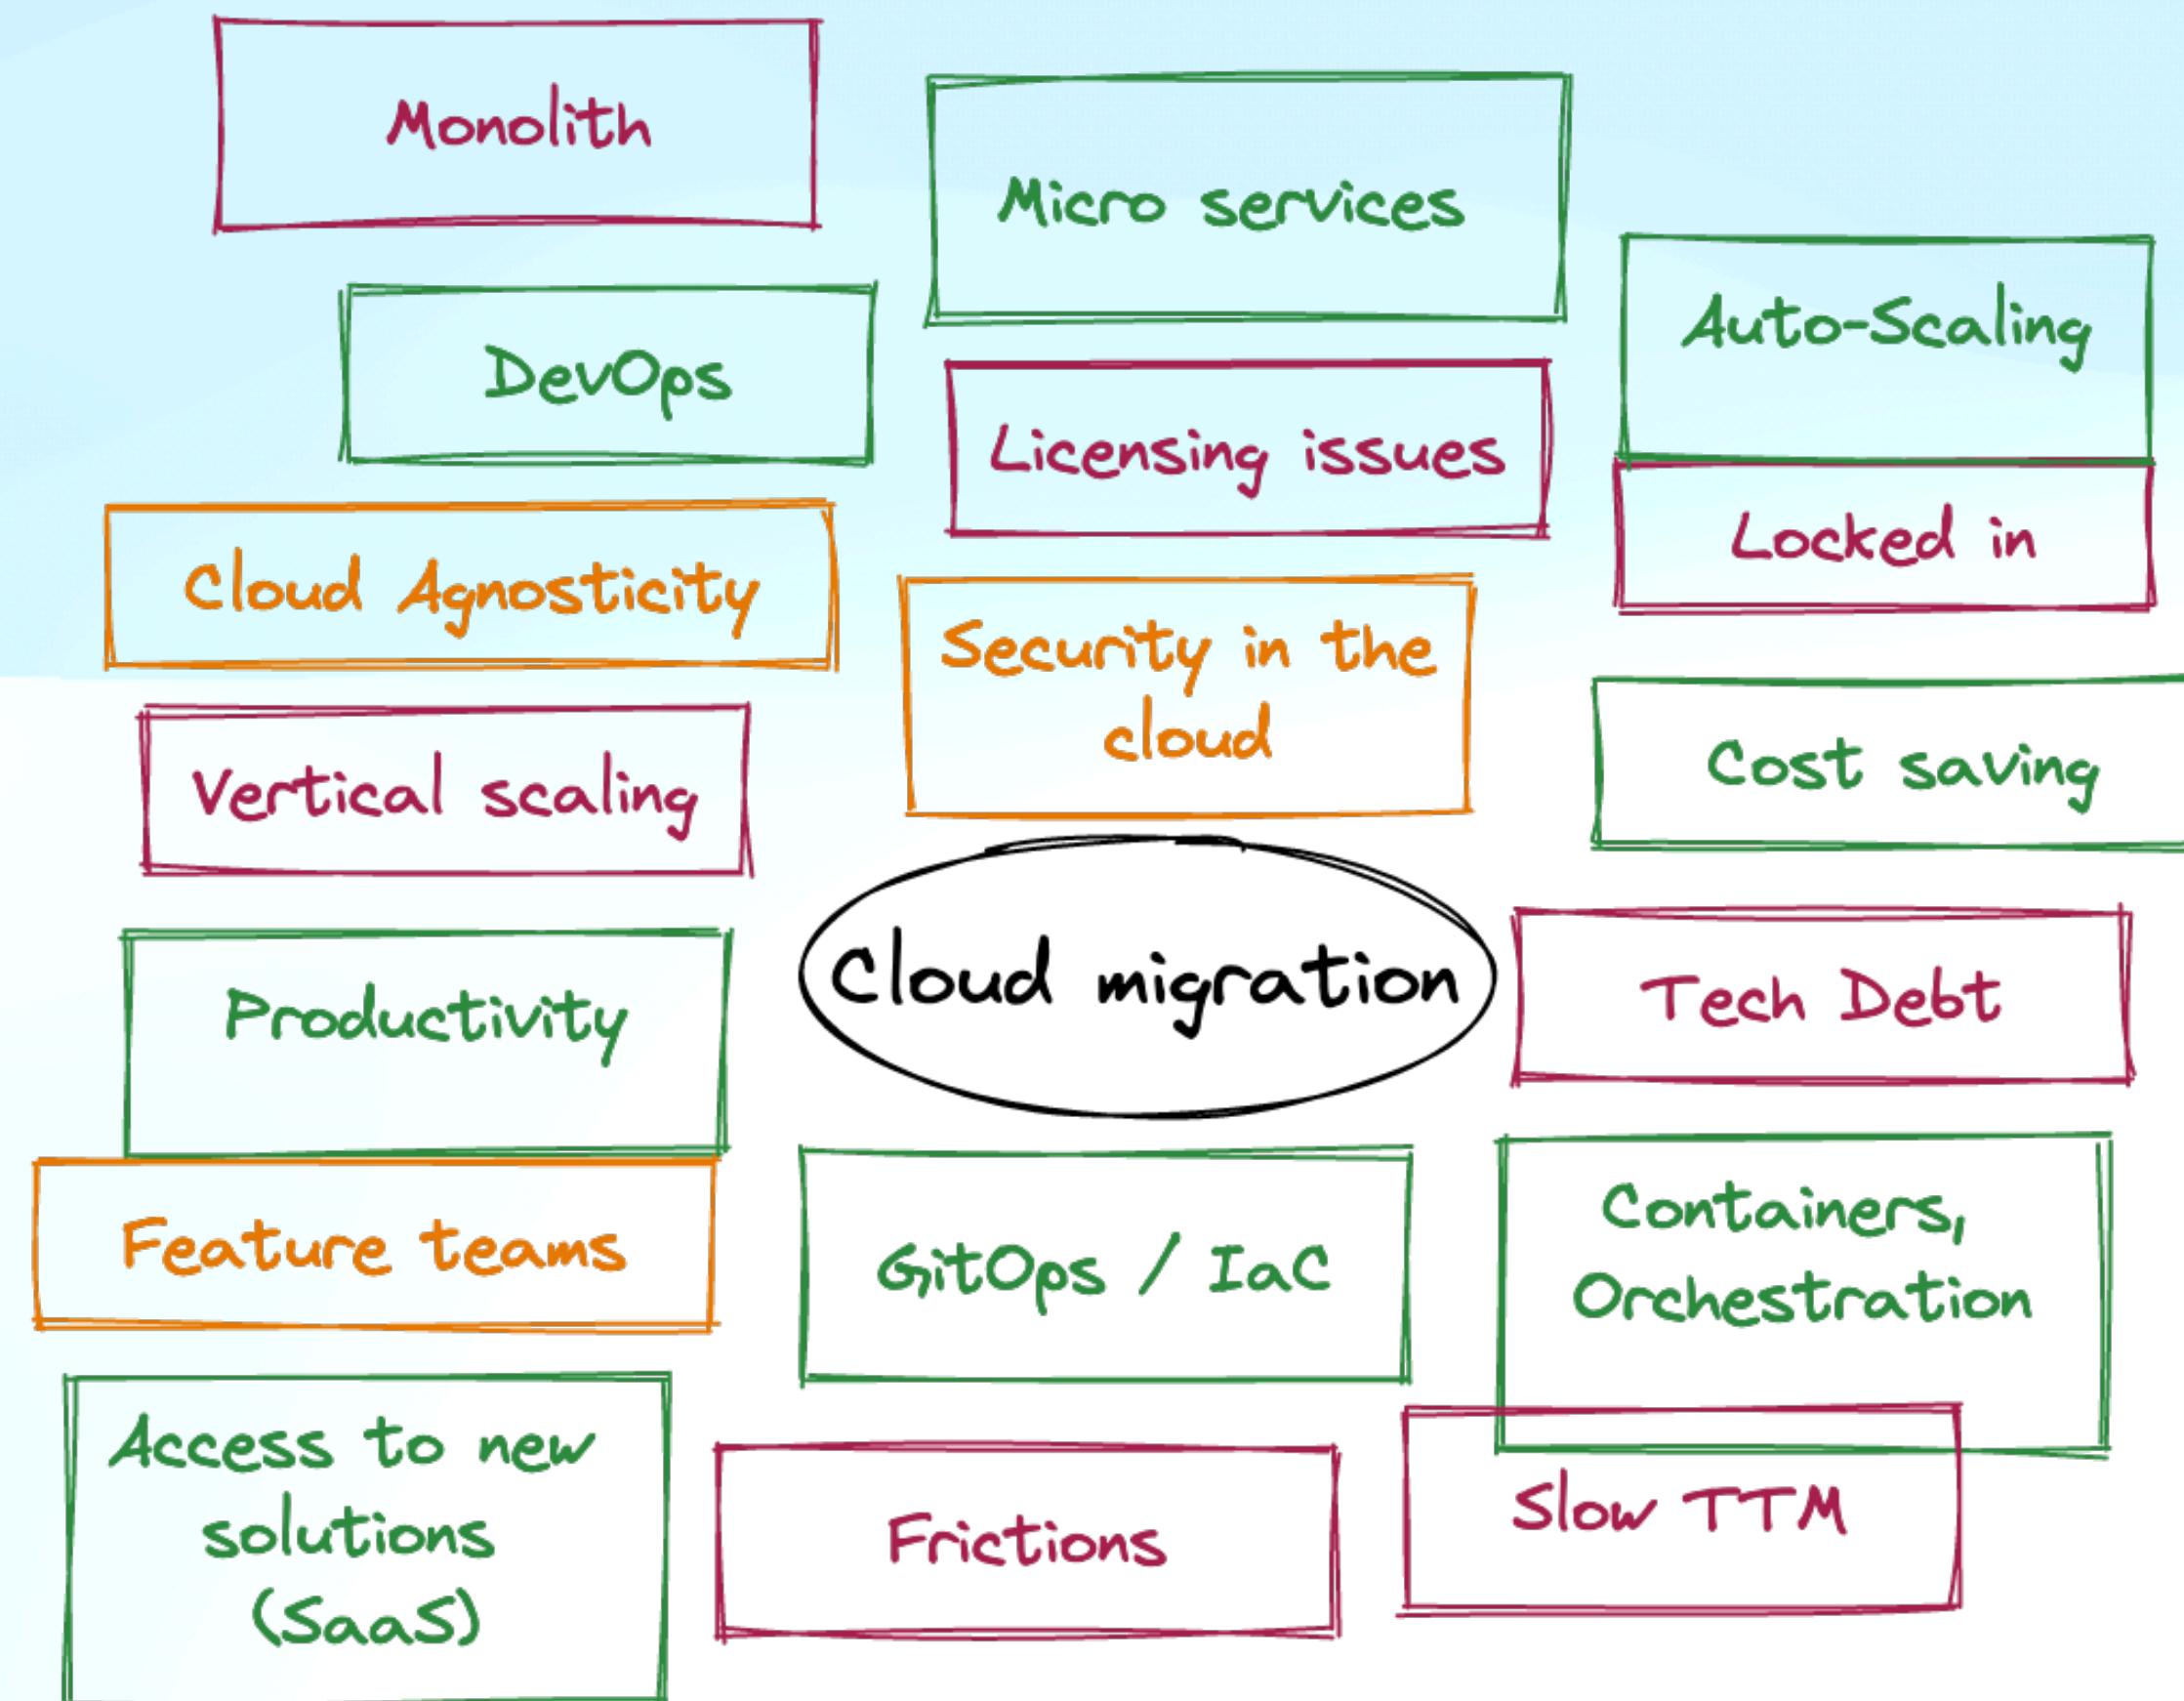
Mapping the problems

The starting line



1. Mapping the problems

Confusion, confusion everywhere



1. Mapping the problems

Analysis Paralysis



Analysis paralysis (or *paralysis by analysis*) describes an individual or group process where overanalyzing or overthinking a situation can cause forward motion or decision-making to become "paralyzed", [...]

A situation may be deemed too complicated and a decision is never made, or made much too late, due to anxiety that a potentially larger problem may arise.

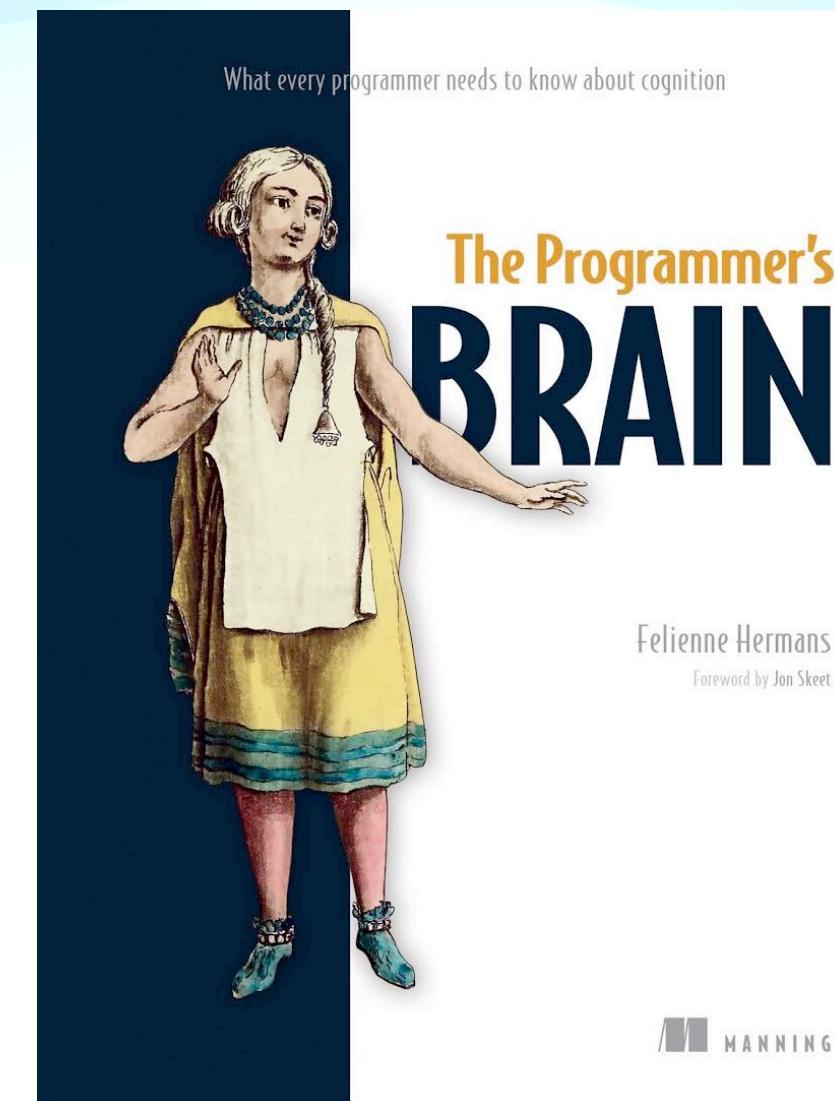
1. Mapping the problems

Identifying the problem

“It’s too big to reason about”

<= Overloaded working memory

- Solutions?
 - Chunk the problem 😕
 - Filter-out noise 😒



👉 Build simpler mental models: “Schemata”

"PAUSE"

Where do we start from?

Tackling Confusion

Defining the schemata



A black hole in space with a blue and purple accretion disk.

Data Gravity

2. Tackling Confusion

DATA GRAVITY

- Where the **data is**, the **services are**,
- Where the **services are**, the **value is**

The **better** the data,

The **better** the service,

The **most** value

2. Tackling Confusion

DATA GRAVITY

- We need to **move the data FIRST**
 - Everything should follow organically
 - The sooner the better
 - In a reliable way
 - No need to “make people migrate”
 - Teams should ask for it

KiSS

2. Tackling Confusion

KiSS

- In case of doubt: **Simpler is better** (less is more)
- Simple != Simplistic, but:
 - **Understand > Implement** (DIY for learning purposes only)
 - No **single solution** to rule them all (Monolith, etc.)
 - Instead:
 - **Single standard / protocol**
 - **Single approach / way of thinking**



Continuously
Better

2. Tackling Confusion

Continuously Better

- It is going to be **long**, accept it
 - No “Big Bang”
 - Unrealistic
 - Too **risky**
 - + Tunnelling effect
 - But **BETTER**:
 - No 1-1 migration => Find **value**



Opportunistic

2. Tackling Confusion

Opportunistic

- No matter what: it **must bring value**
 - Avoid the “tech for techies”
- Exercises:
 - Police lineup (usual suspect)
 - Or the fridge analogy (eat what you have)
- Embrace opportunities
 - You know them, you’ll find them along the way

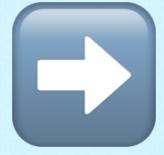
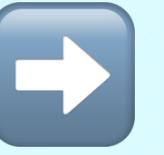
Dogfooding



2. Tackling Confusion

Dogfooding

- “Are we doing this right”

 **Most FAQ**  **Mental Charge**

- A way to build convictions:

- “Does this work for us?”
 - “Do we trust ourselves as a provider?”
- Quality matters:
 - Build exigence
 - (Internal+External) Audit



**Decisions
must NOT matter**

2. Tackling Confusion

Decisions must not matter

- Standards, standards, standards
- OSS
- Documentation (transparent) > implementation
 - ADRs
 - ...
- Favour *reversible* / *agnostic* when possible

=> piggy-backed in KiSS

Our schemata

2. Tackling Confusion

Our pillars

- Move the data (Data Gravity)
- Only the reliable part of it, cleaned (Better)
- Using standards, in a simple way (KISS)
- Progressively (Continuous)
- not sure where we'll store it (Decisions / Opportunities)
- but we'll use it ourselves (Dogfooding)

2. Tackling Confusion

Build your own decision framework

- Your own schemata
- Whatever helps your teams

Ready to screen the migration candidates?

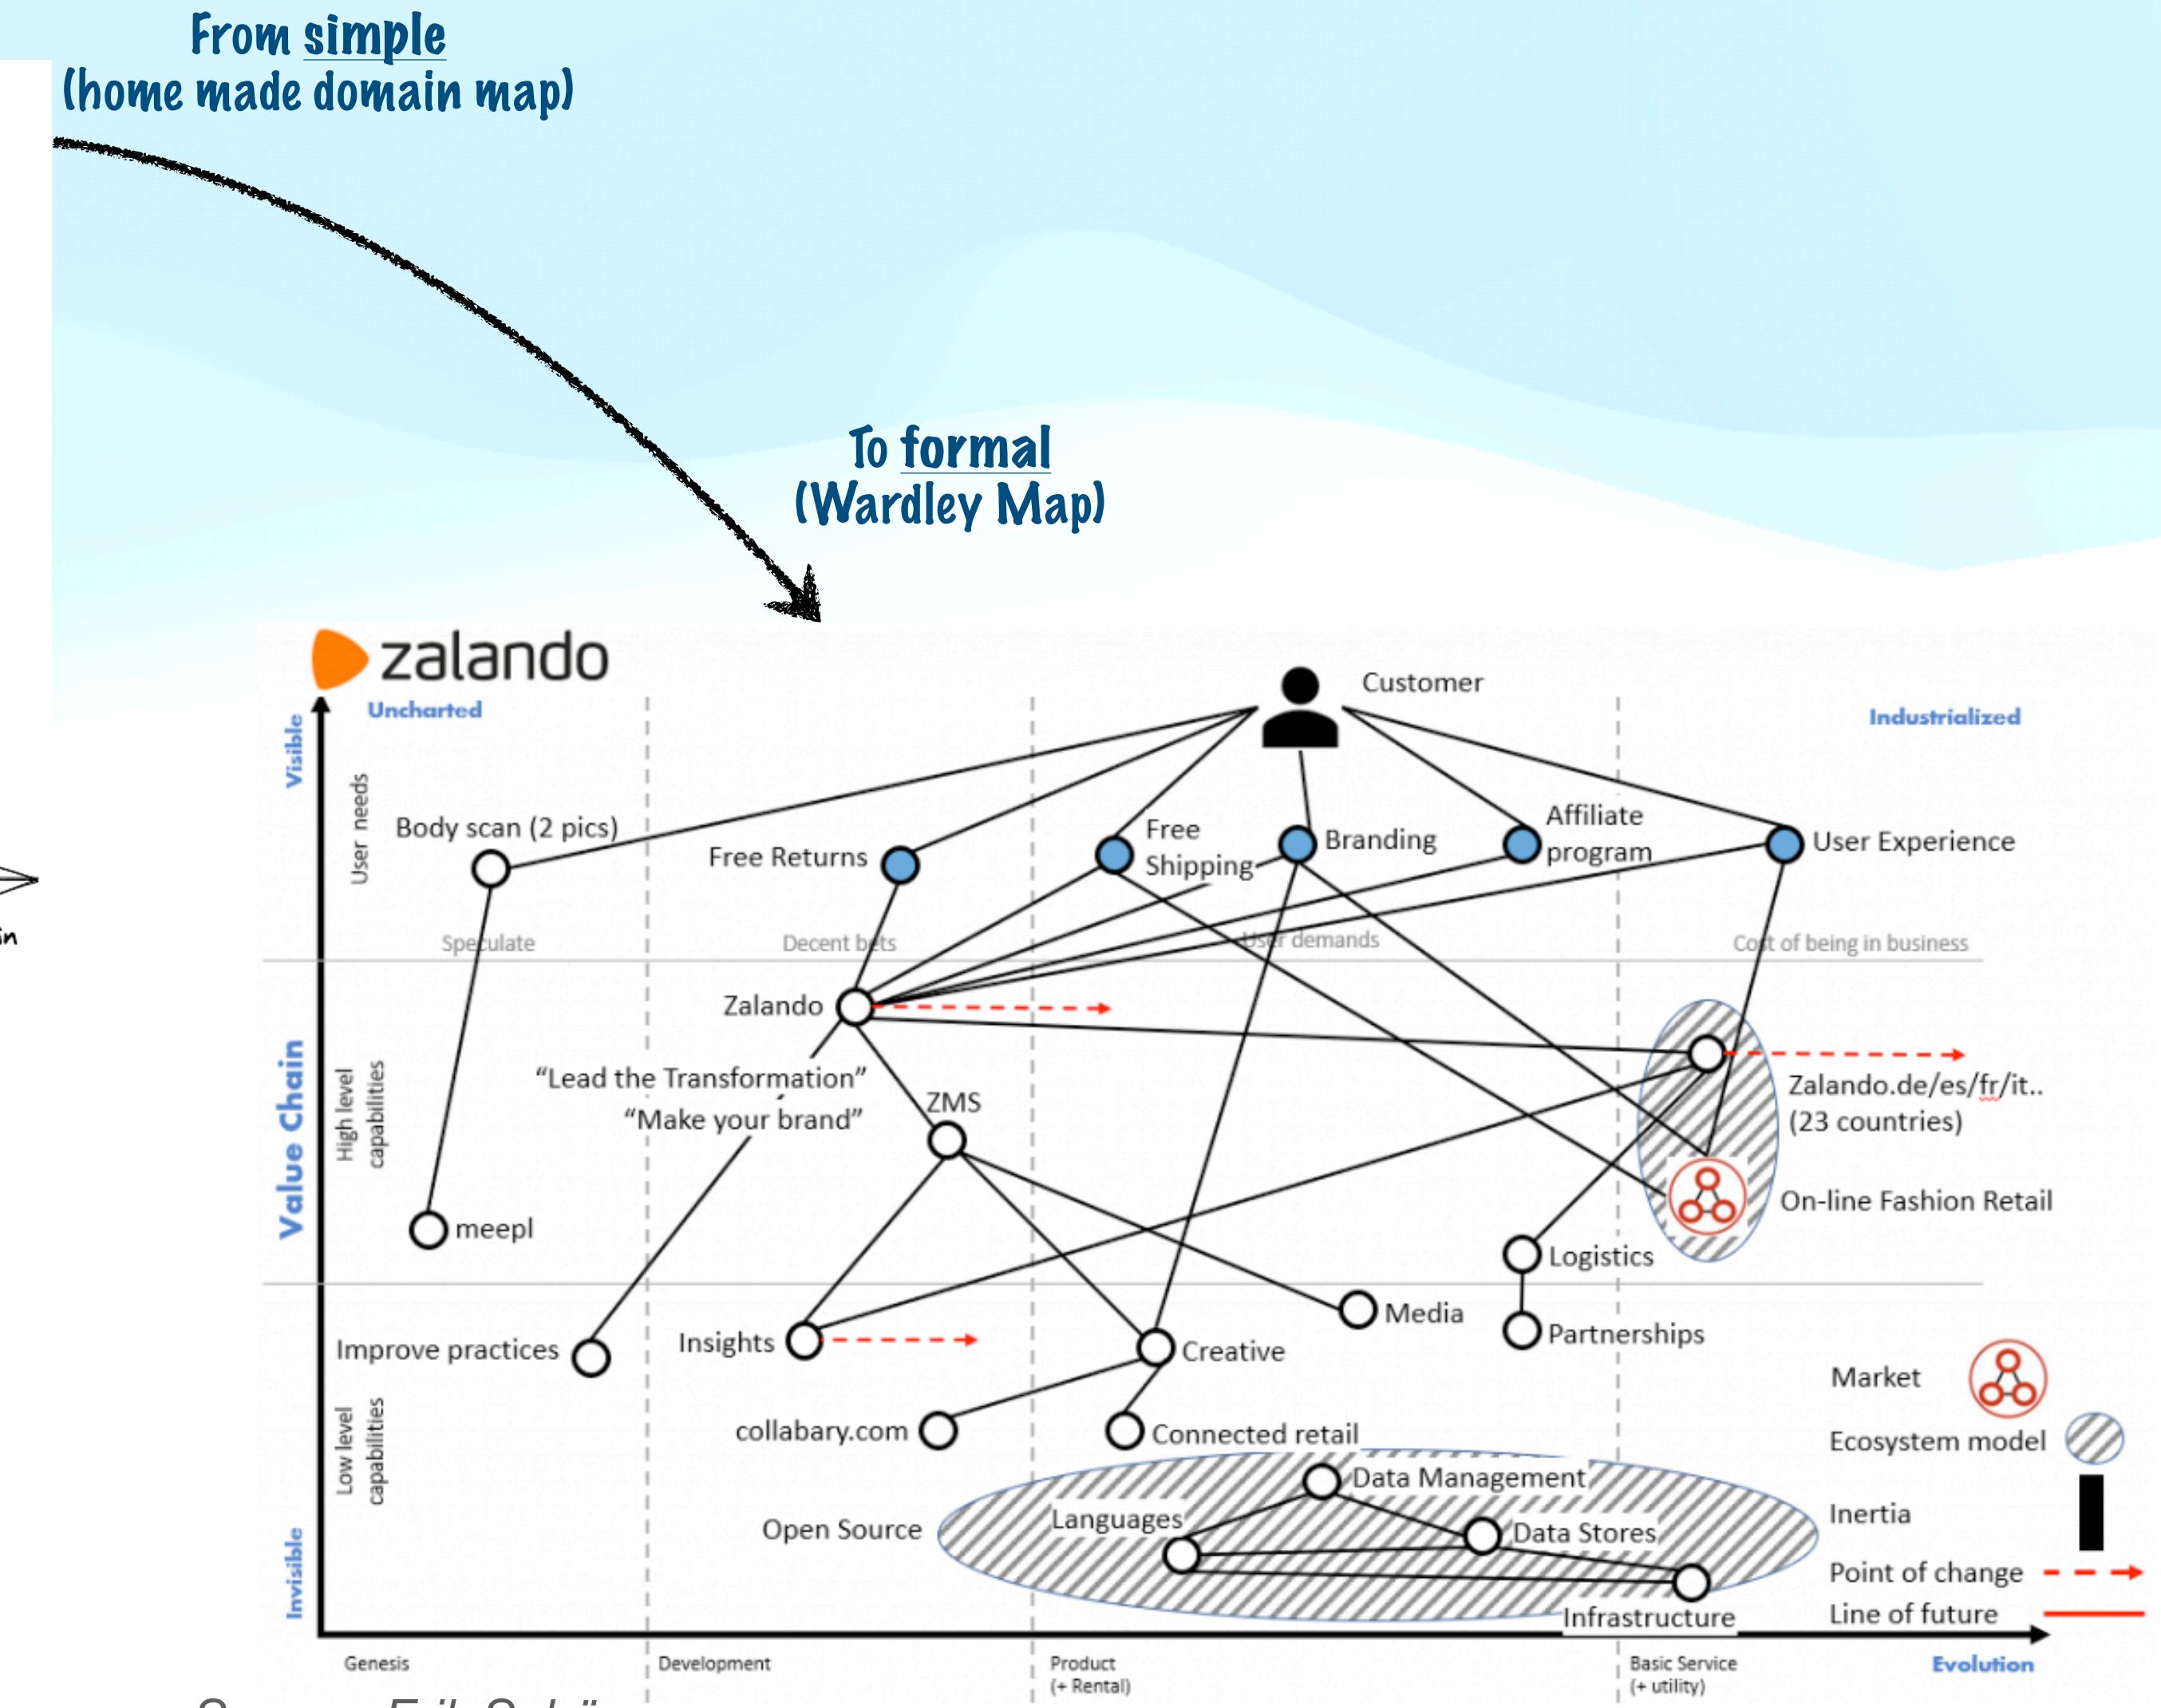
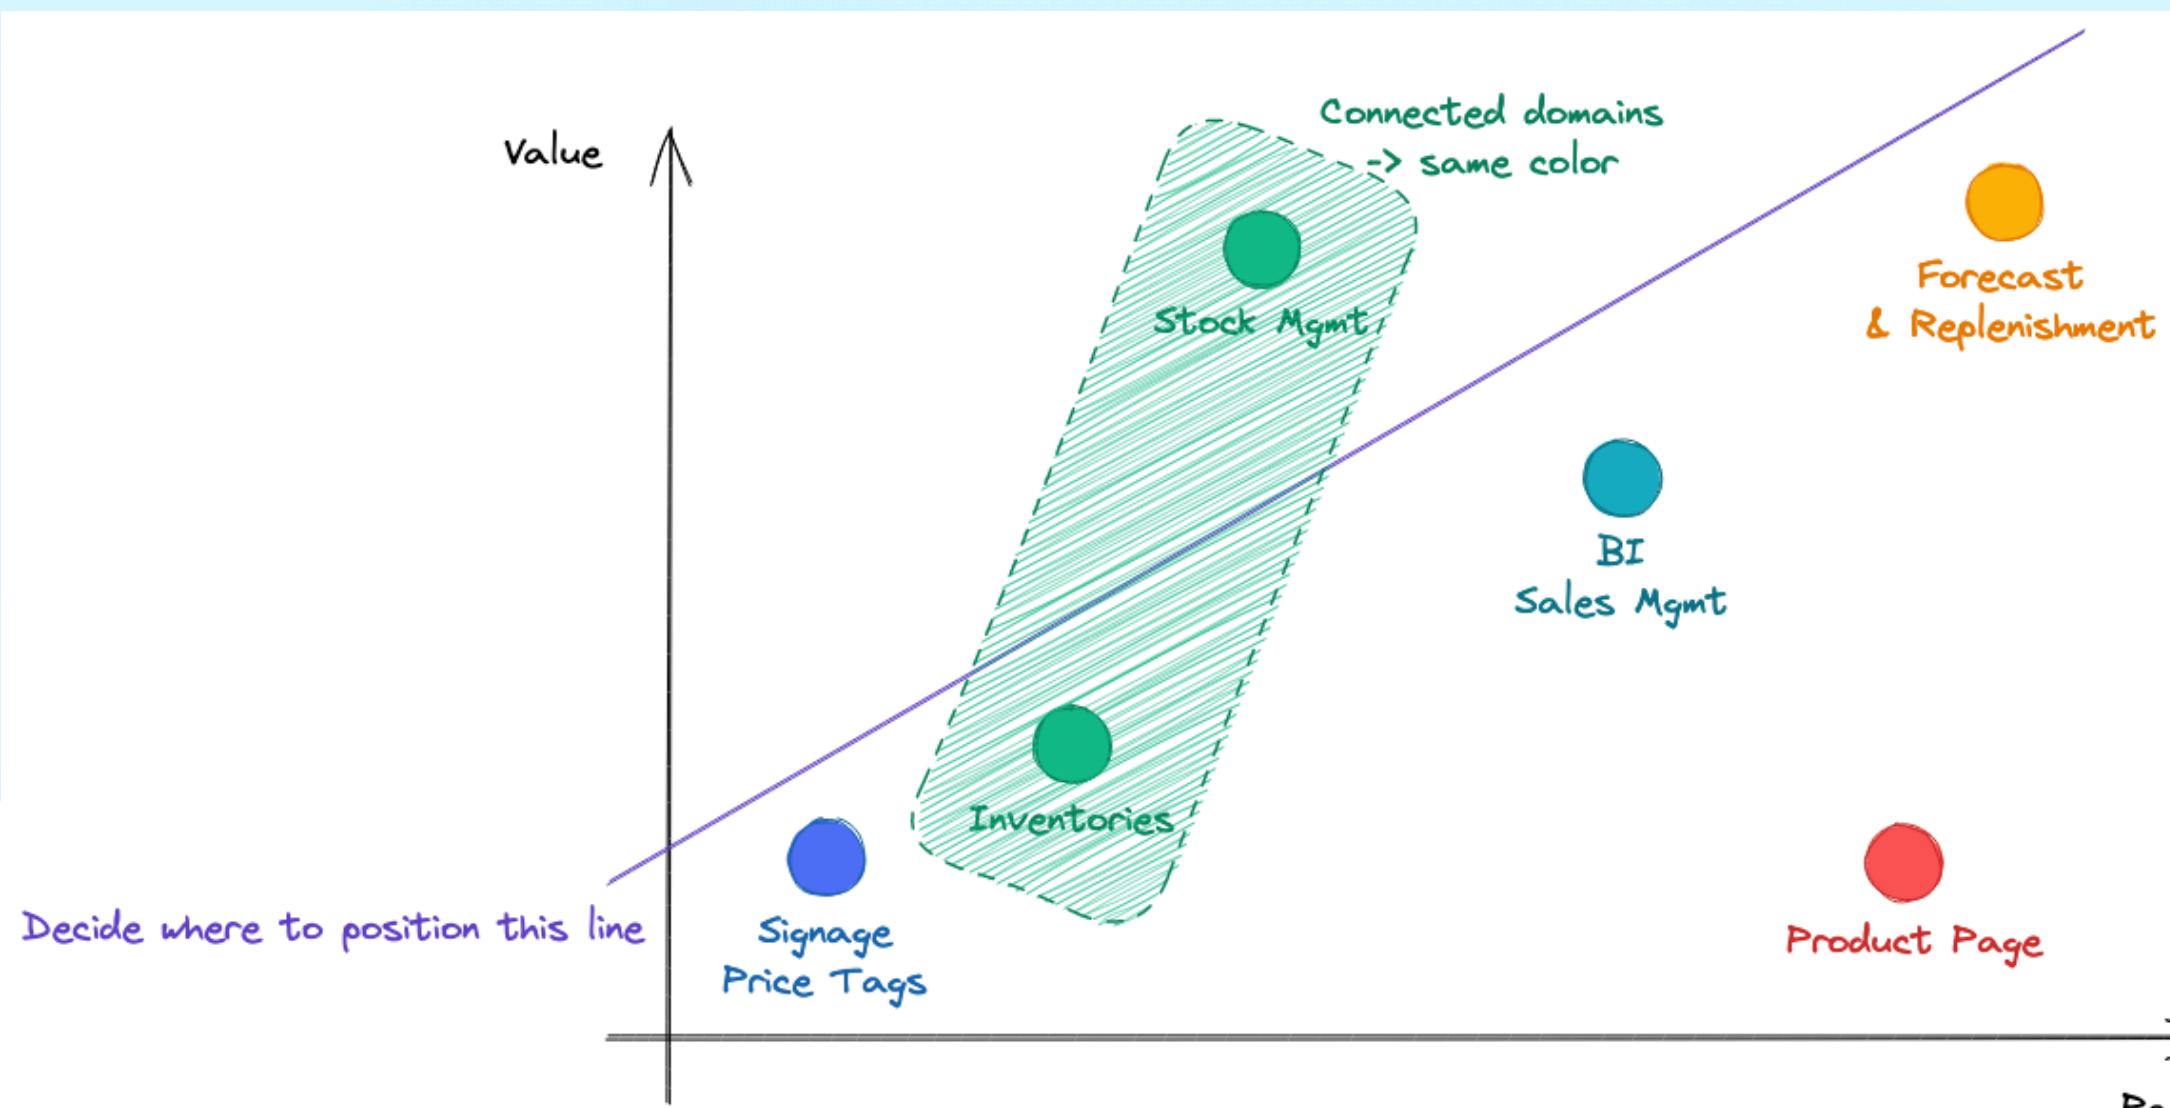
2. Tackling Confusion

WHAT WE TRIED



2. Tackling Confusion

WHAT WORKS FOR YOU



"PAUSE"

Where do we start from?

*Move data,
In the simplest way,
Continuously, better,
Bringing value along the way,
For us first,
=> **Domain***

What do we start with?

Event Streaming

The tactics



3. Event Streaming

Move what exactly?

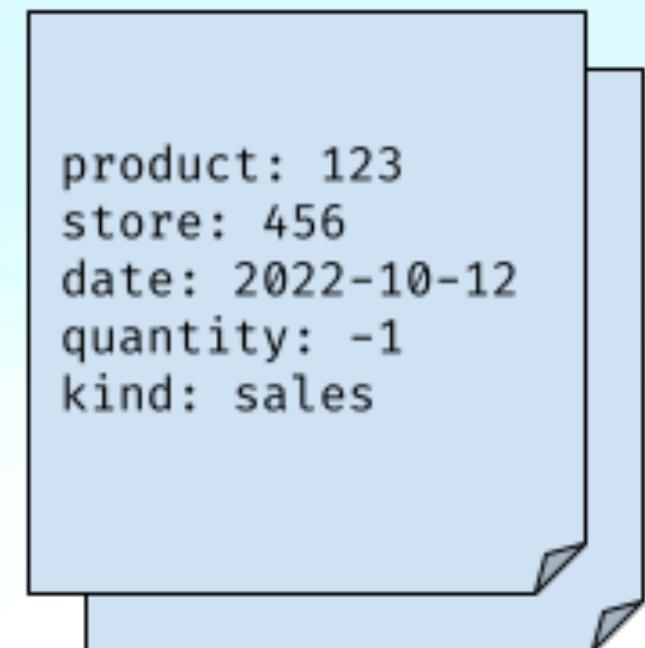
- Model#1: “Move the data”
 - ➡ Which data?
- Identified domain, but broadly
 - eg: Stocks, what stocks?

3. Event Streaming

Streaming what?

Stock Management example

Movement

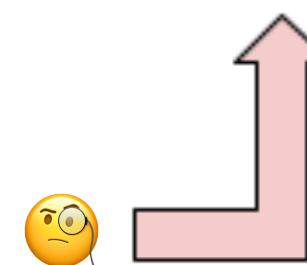


Few usage (mostly internal)
Displayed in the UI
Flowed to supply teams
No real time use

Picture

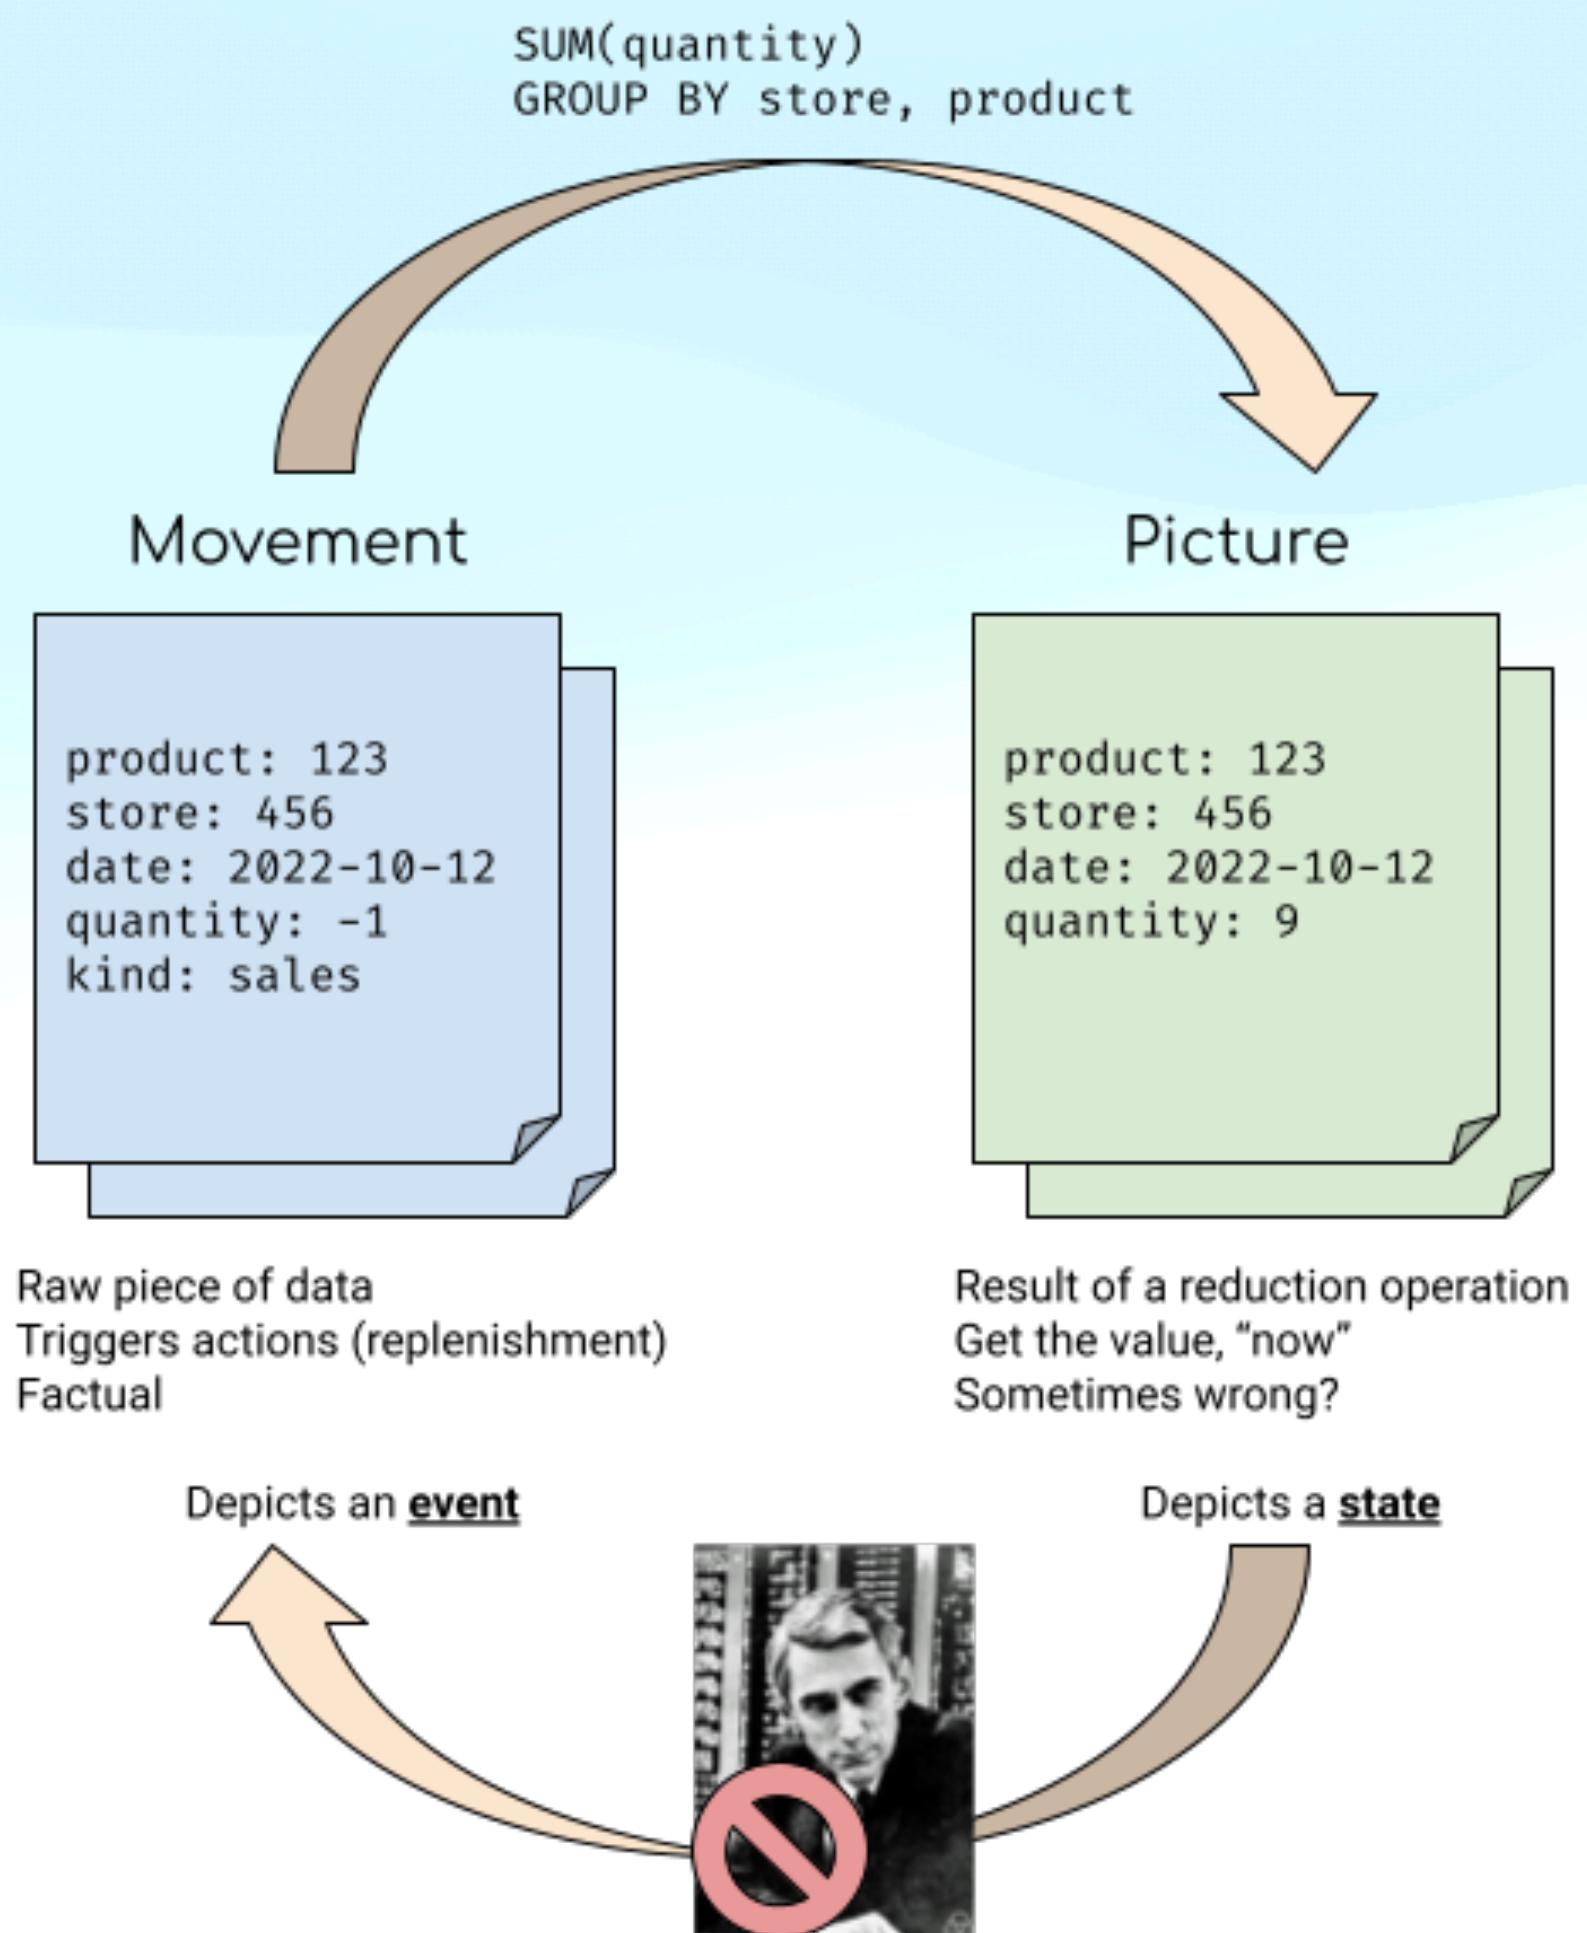


Most used API in the company
Everyone needs it (e-comm, etc)
Highly critical
High source of problems



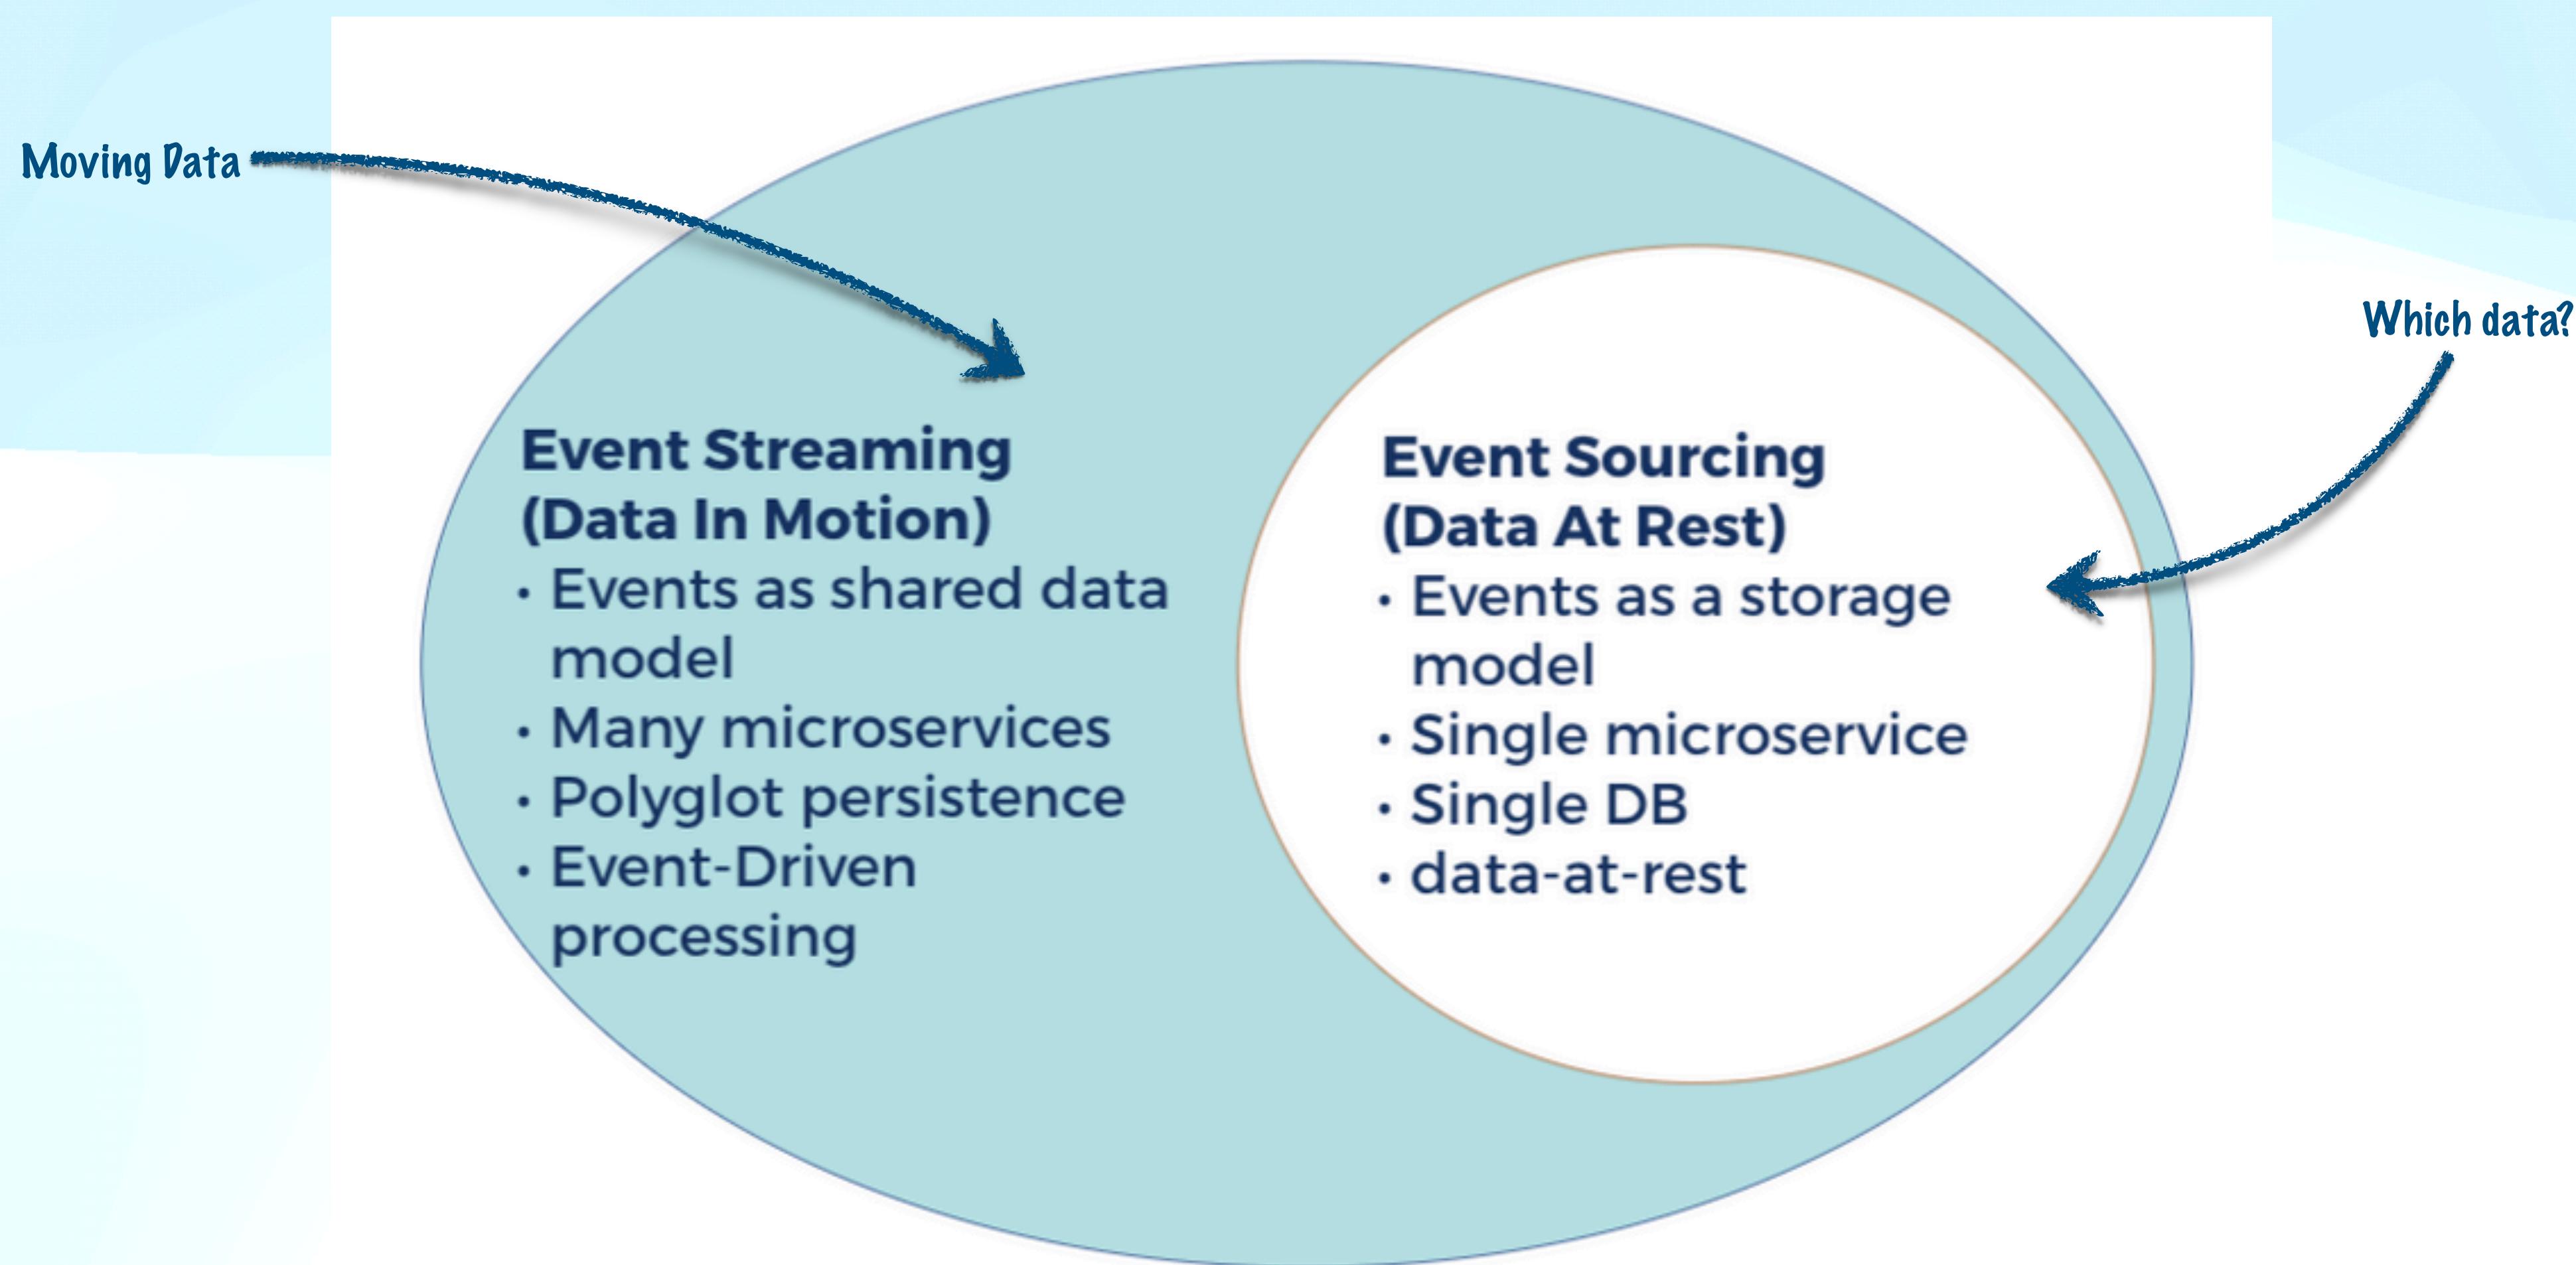
3. Event Streaming

Getting to the source of truth



3. Event Streaming

Continuous = Streaming / What = Event



3. Event Streaming

A common pattern, across common industries

- Events are everywhere:
 - Financial transactions
 - IoT sensors
 - Traceability
 - etc. => Choose your **DOMAIN EVENTS**
 - (and remember Claude)

3. Event Streaming

From **means** to **end**

- Teams loved it
 - Build a local cache (**Continuously, Better**)
 - Notification systems (**Opportunities**)
 - “Let us deal with this” == “Send us the event” (Chunking => KiSS)
 - Become a generic event provider => simpler scope (KiSS)
- Customers loved it (**Dogfooding**)
 - ~~Arbitrary delays~~ => Events in real time
 - “Tell me what’s happening” / “Notify me”

"PAUSE"

Where do we start from?

*Move data,
In the simplest way,
Continuously, better,
Bringing value along the way,
For us first,
=> A domain*

What do we start with?

*Streaming Events,
Shaping a source of truth*

How do we do that?

Picking the right Tech

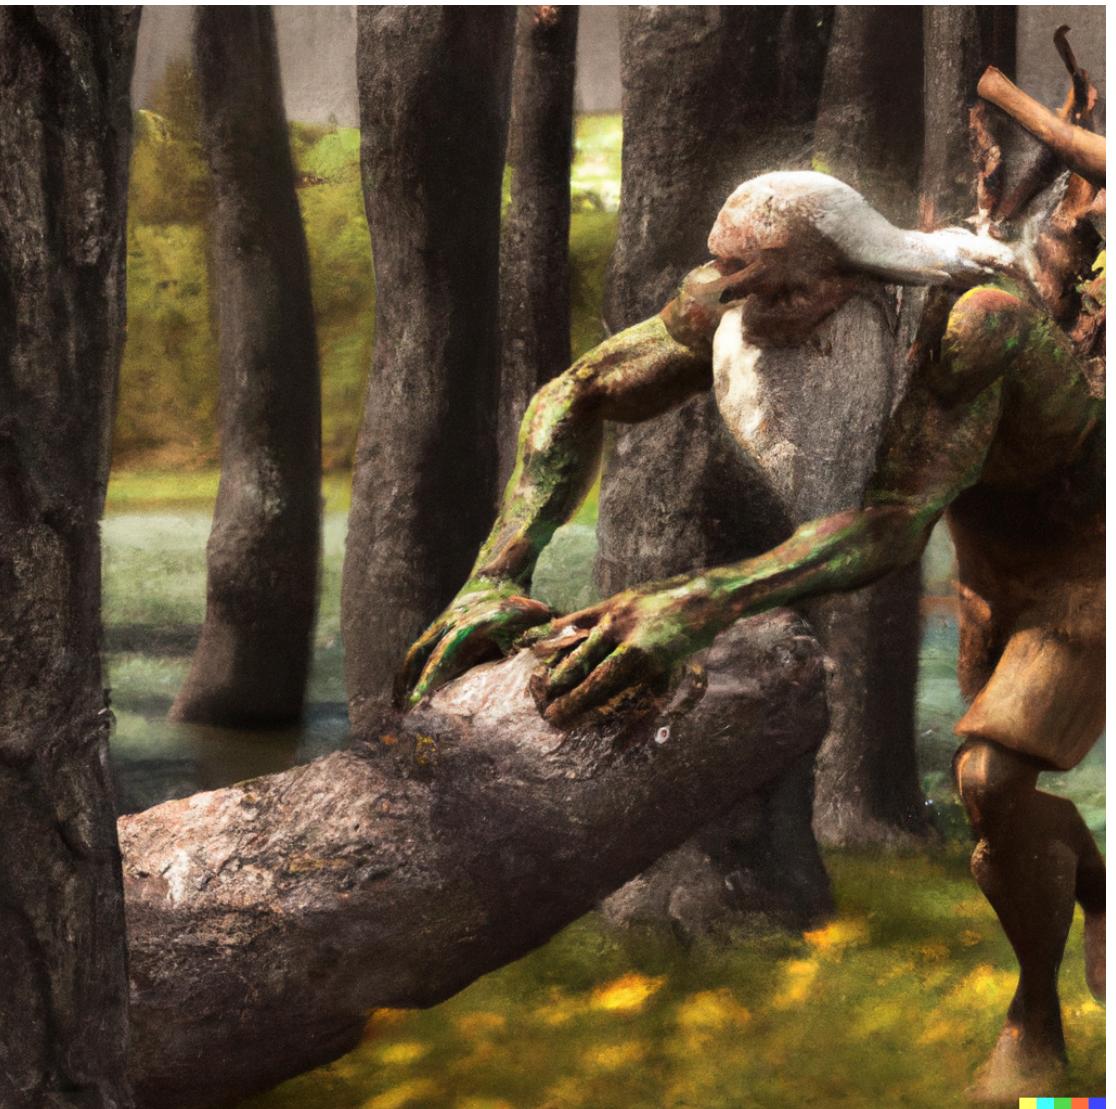
Apache Kafka



4. Picking the right tech

Does Kafka help streaming events? 🤔

- Events can be represented as logs
 - Factual <=> Immutable
 - Chronological <=> Ordered
 - Index <=> Offset
 - Bookmark <=> Consumer offset



4. Picking the right tech

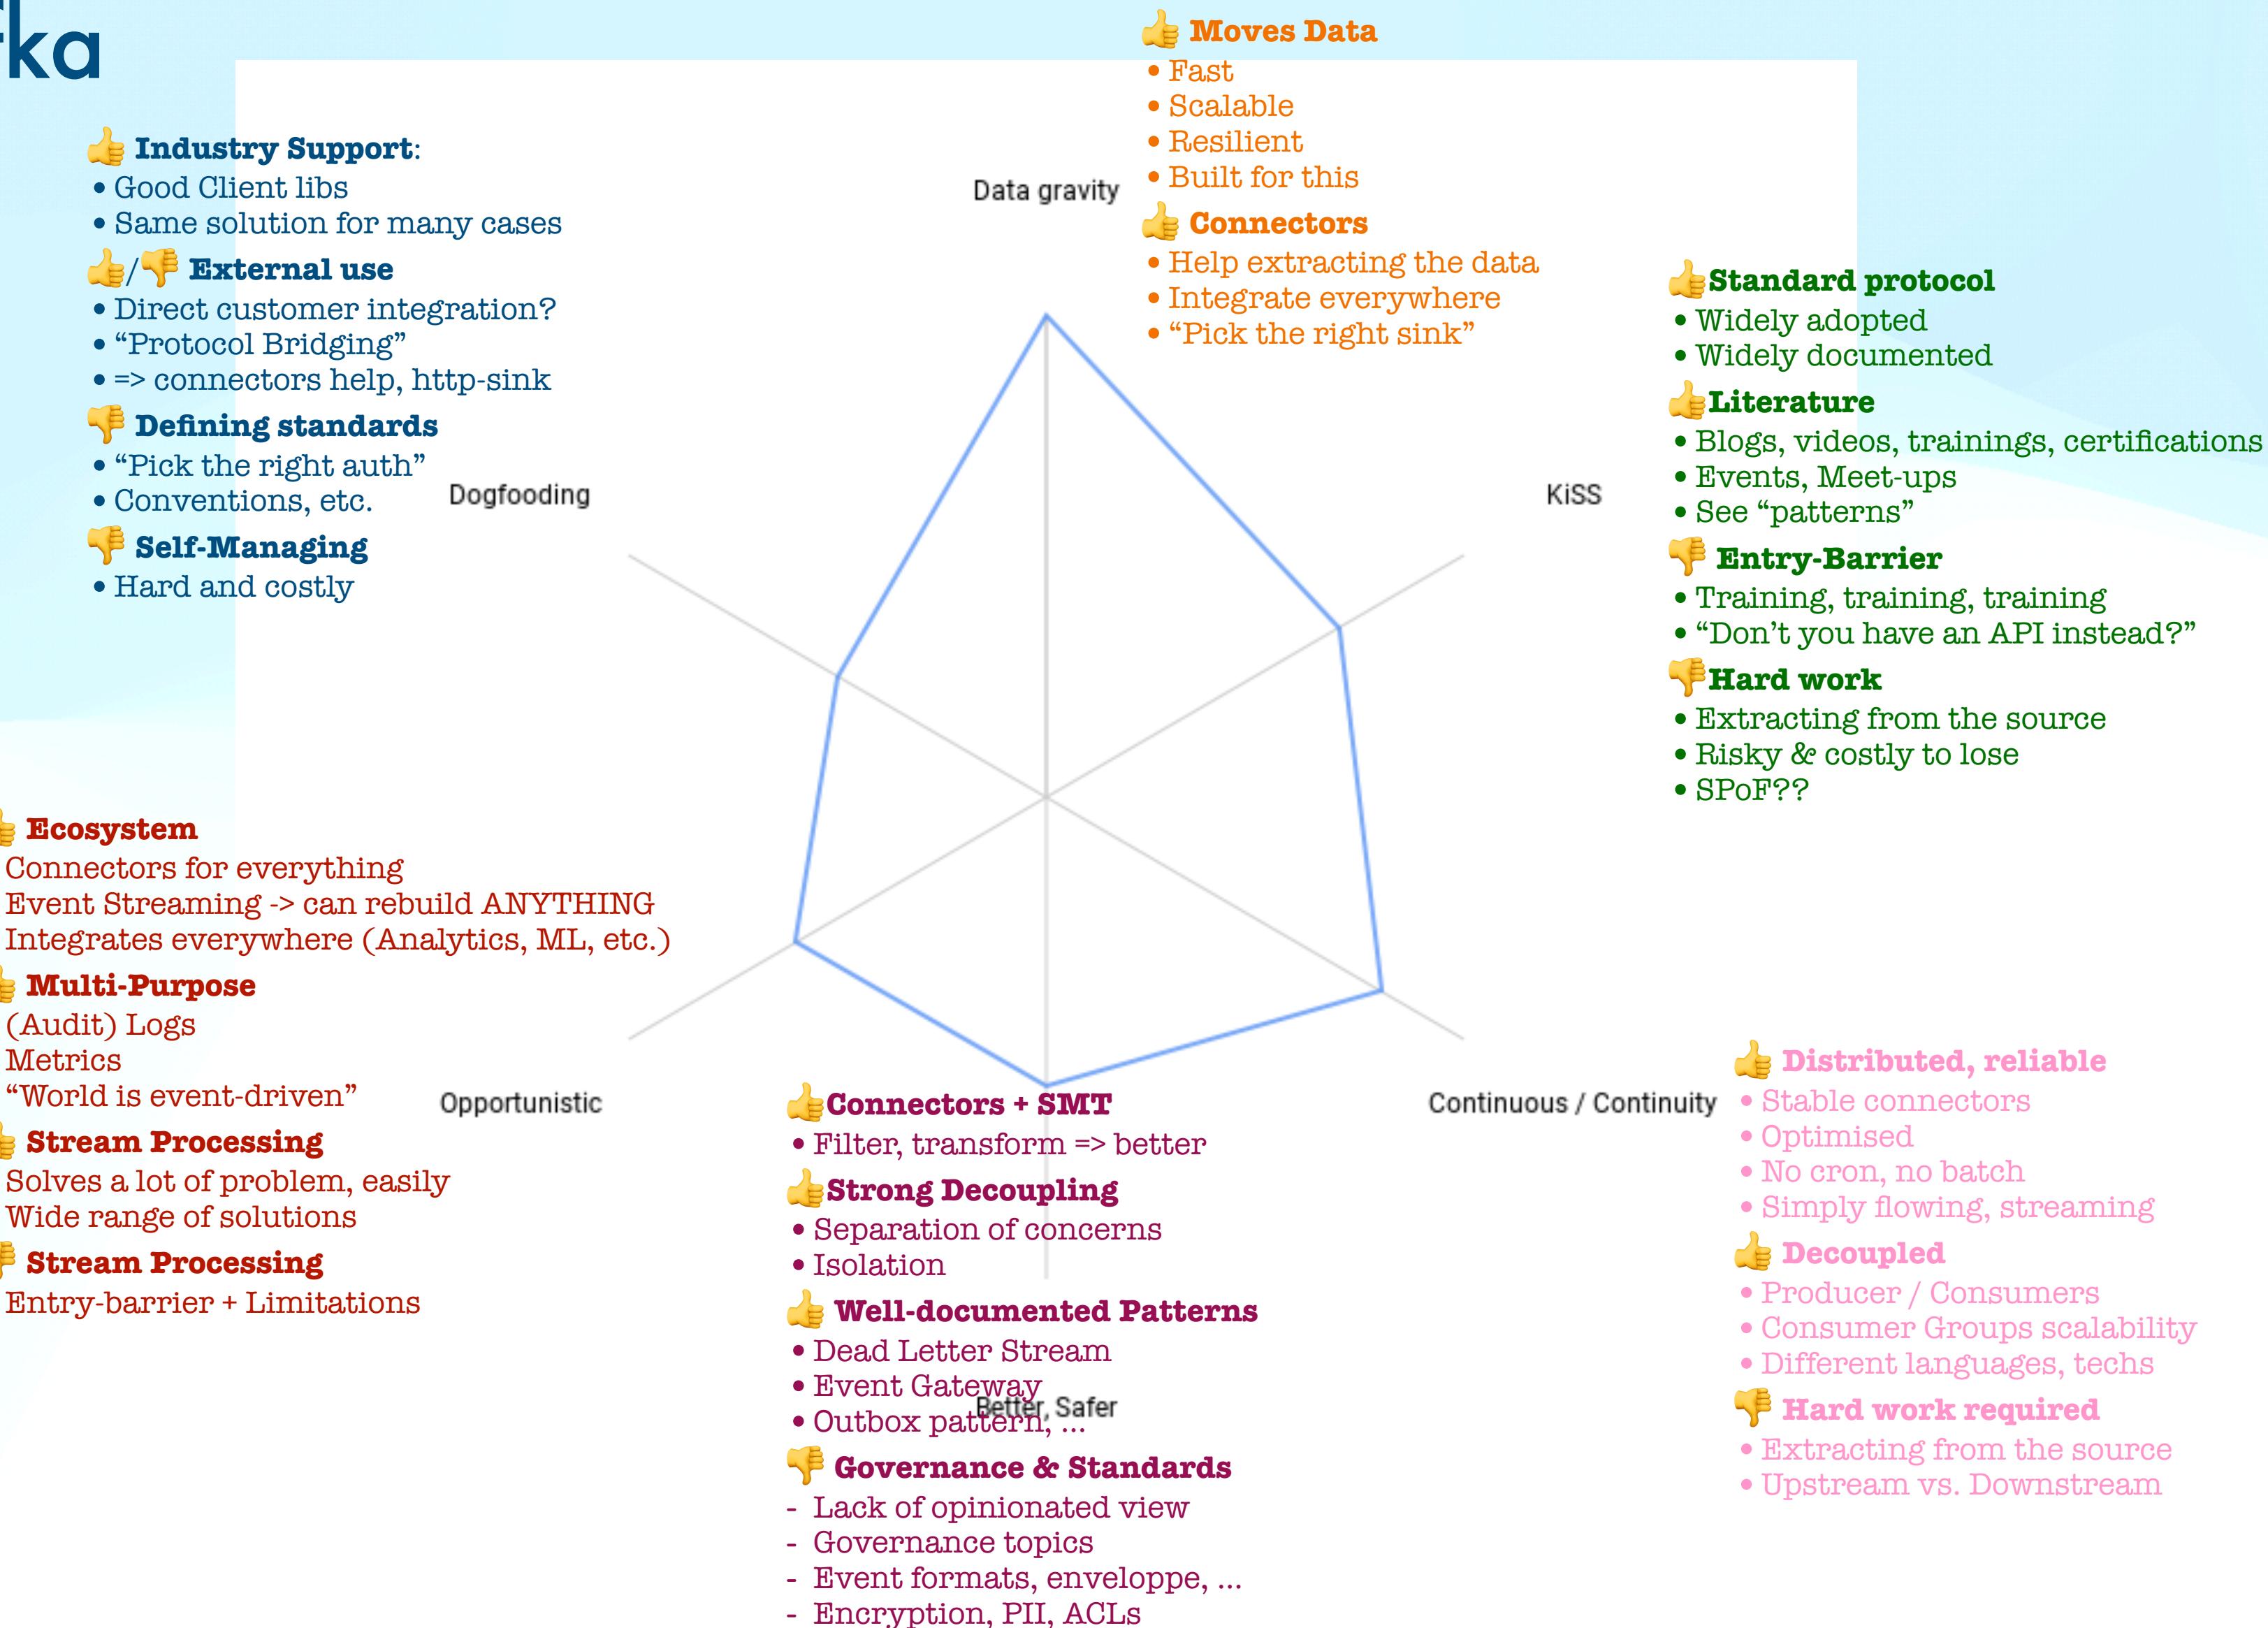
Key features in the context of a migration

- Replay-ability is key (Better, Safer)
- AK (as a protocol) is a standard (Decisions -> KiSS)
- Suits other patterns as well: CQRS, ... (KiSS)
- Vibrant ecosystem (Data Gravity, Better, Decisions -> KiSS)
- AK is not simple, but well documented (Decisions -> KiSS)
- Stream Processing (Better, Opportunities)
- Advanced Governance features (Better, Safer)



4. Picking the right tech

Evaluating Kafka



"PAUSE"

Where do we start from?

Move data,
In the simplest way,
Continuously, better,
Bringing value along the way,
For us first,
=> **A domain**

What do we start with?

Streaming Events,
Shaping a **source of truth**
That can be used by many

How do we do that?

Using Apache Kafka
AND ITS ECOSYSTEM

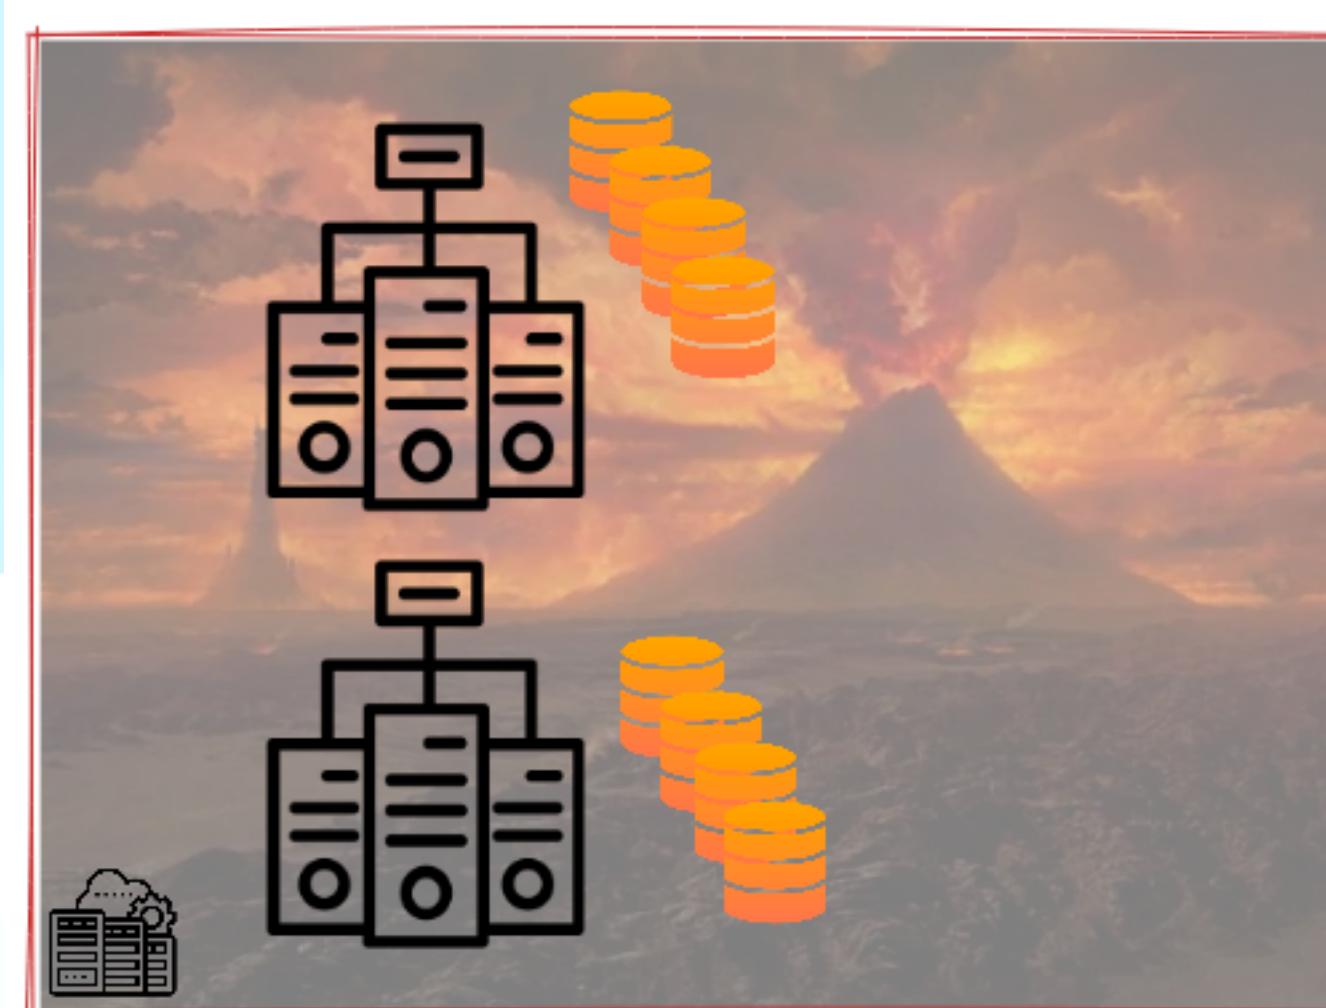
Lessons Learned

And what's next?



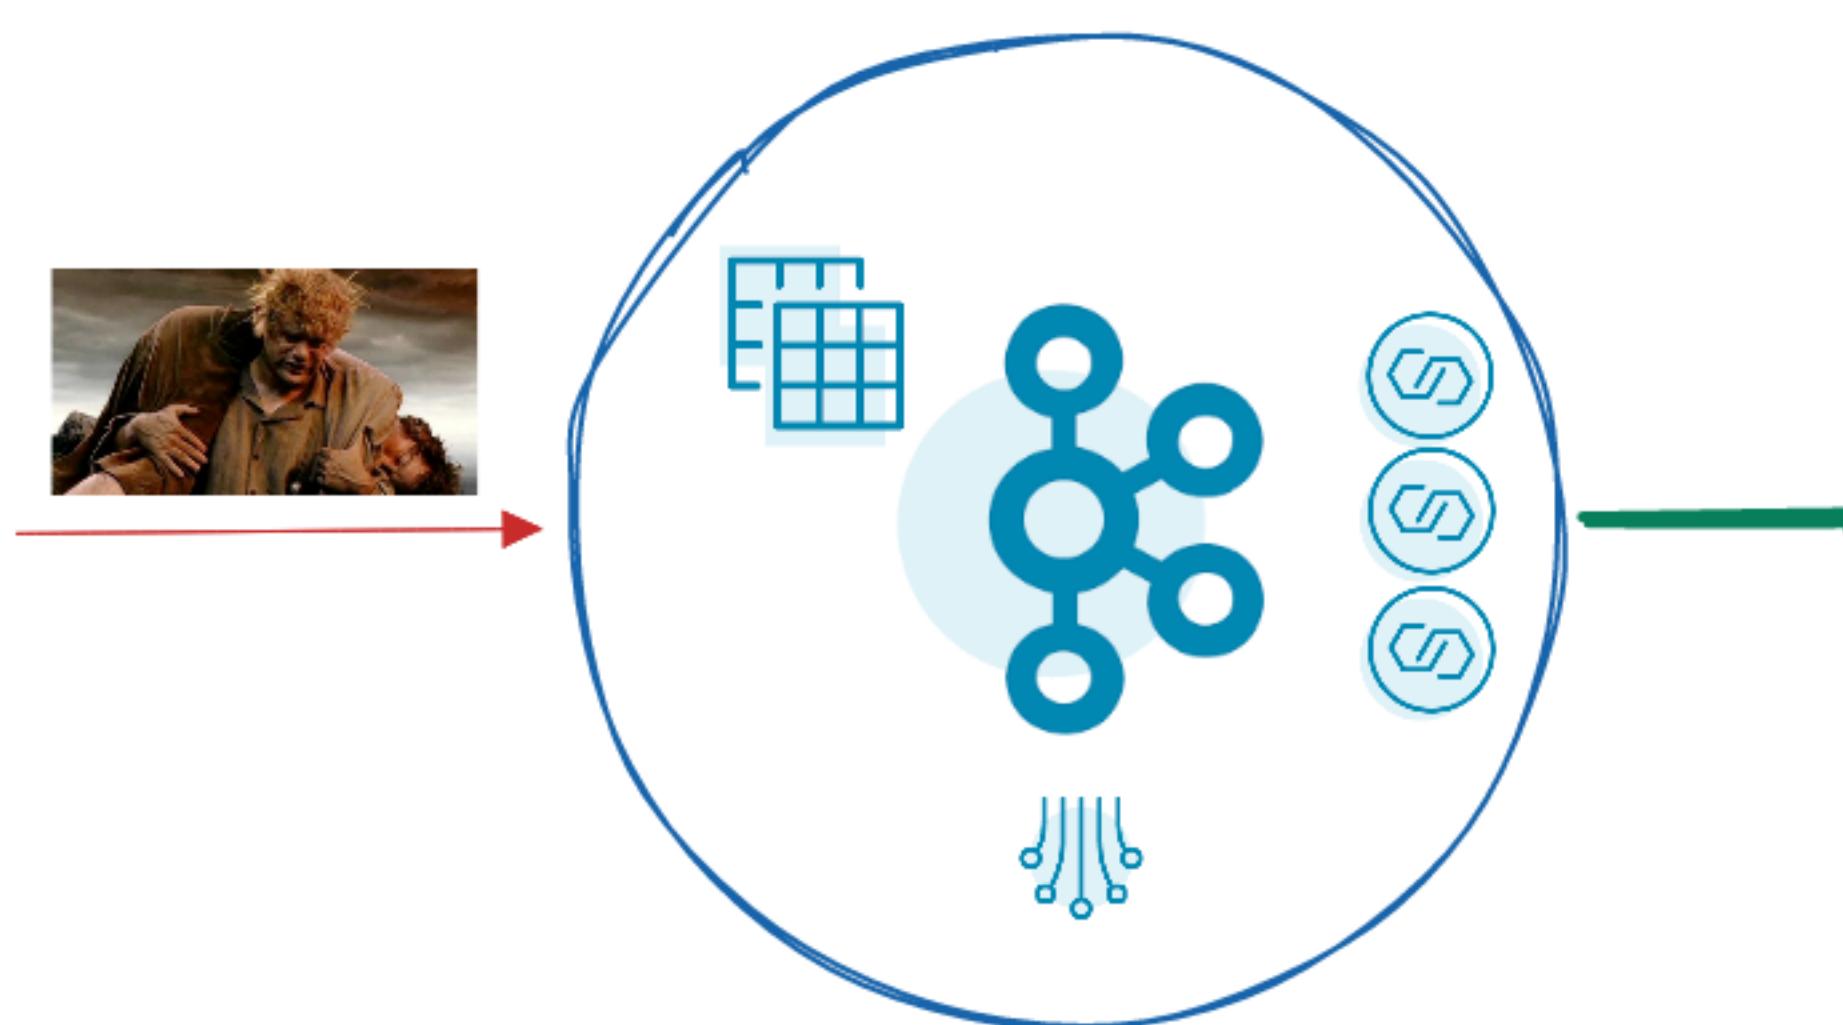
5. Lessons Learned

👍👎 A divided world



Focus efforts upstream
KiSS is strategic here
“Temporary” is acceptable

Move the data, at **all** costs



Low effort (connectors, etc.)
“Dogfooding” is strategic here
“Temporary” is unacceptable

Use the data, at **low** cost

5. Lessons Learned

👍 Reduced Friction

- An answer to a lot of problems
 - compacted topics, stream processing, connectors, etc.
- Isolation of concerns
- Examples:
 - “Can you change your API?” <= most common question ever
 - “What about you consume the data?”

5. Lessons Learned

👎 It's not that easy

- Lots of standards, lack of dogma
 - AuthN / AuthZ
 - Hard to make decisions
- Traps (See “AK: Tips and Tricks I wish I had known before”)
- Self-Maintaining Distributed systems is hard



5. Lessons Learned

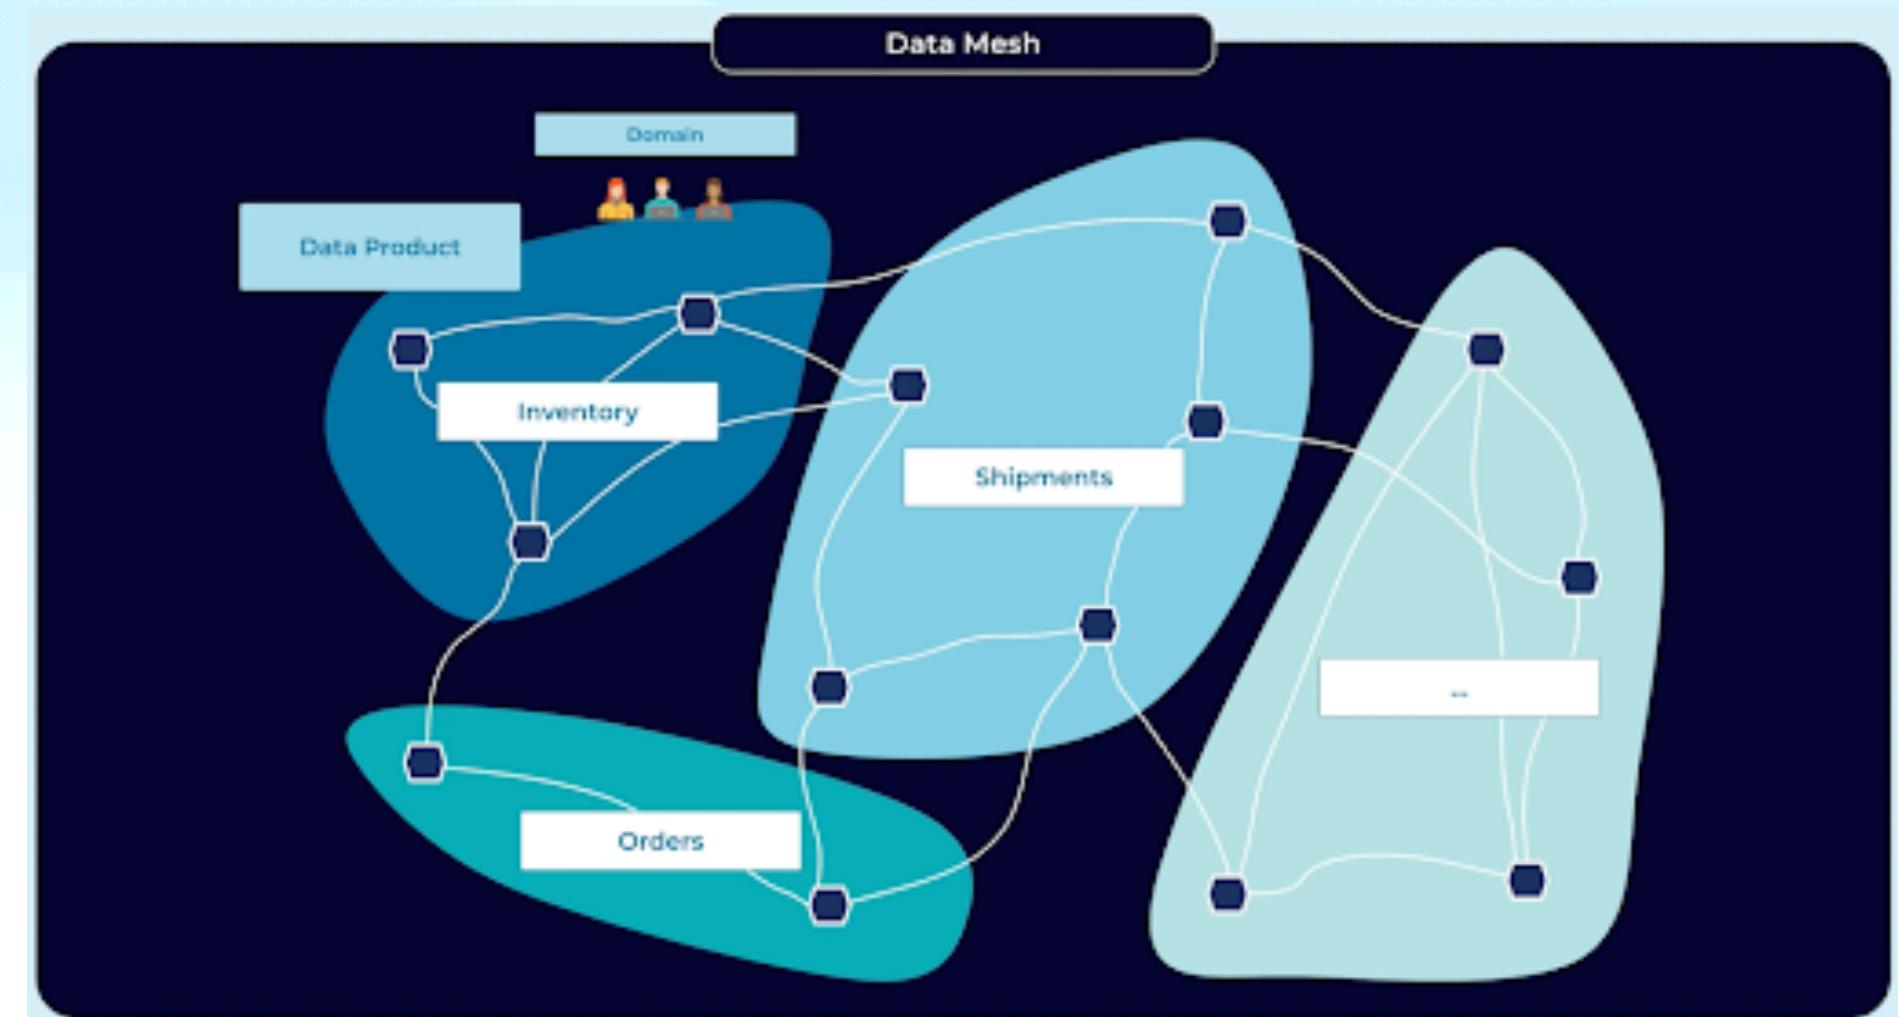
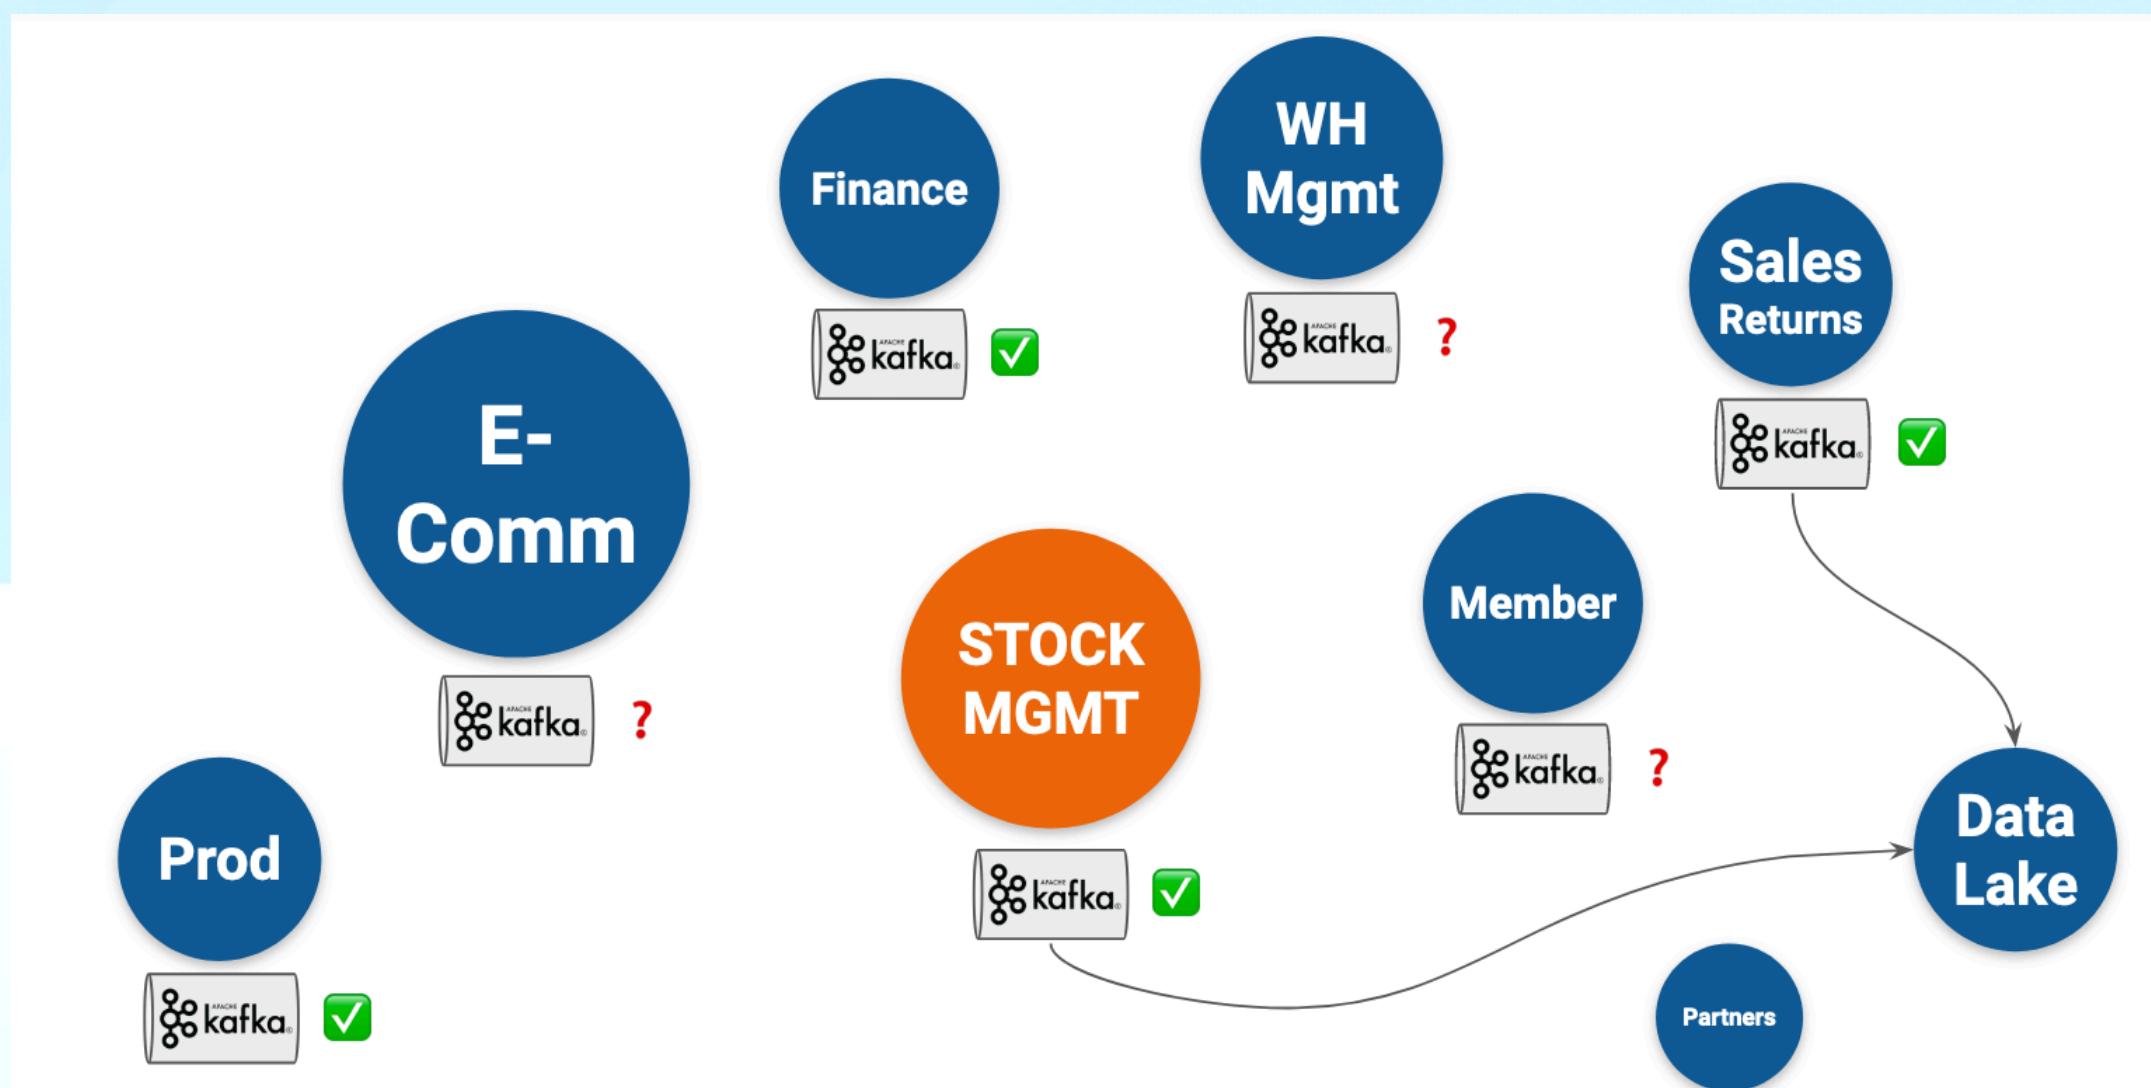
👎 Many hard decisions at scale

- Internal Standards
 - Naming conventions,
 - Data format / schema, ...
- Internal Guidance / Training / Knowledge
 - Dimensioning / Sizing (partition count, clusters)
 - **Document** the traps faced, errors made



5. Lessons Learned

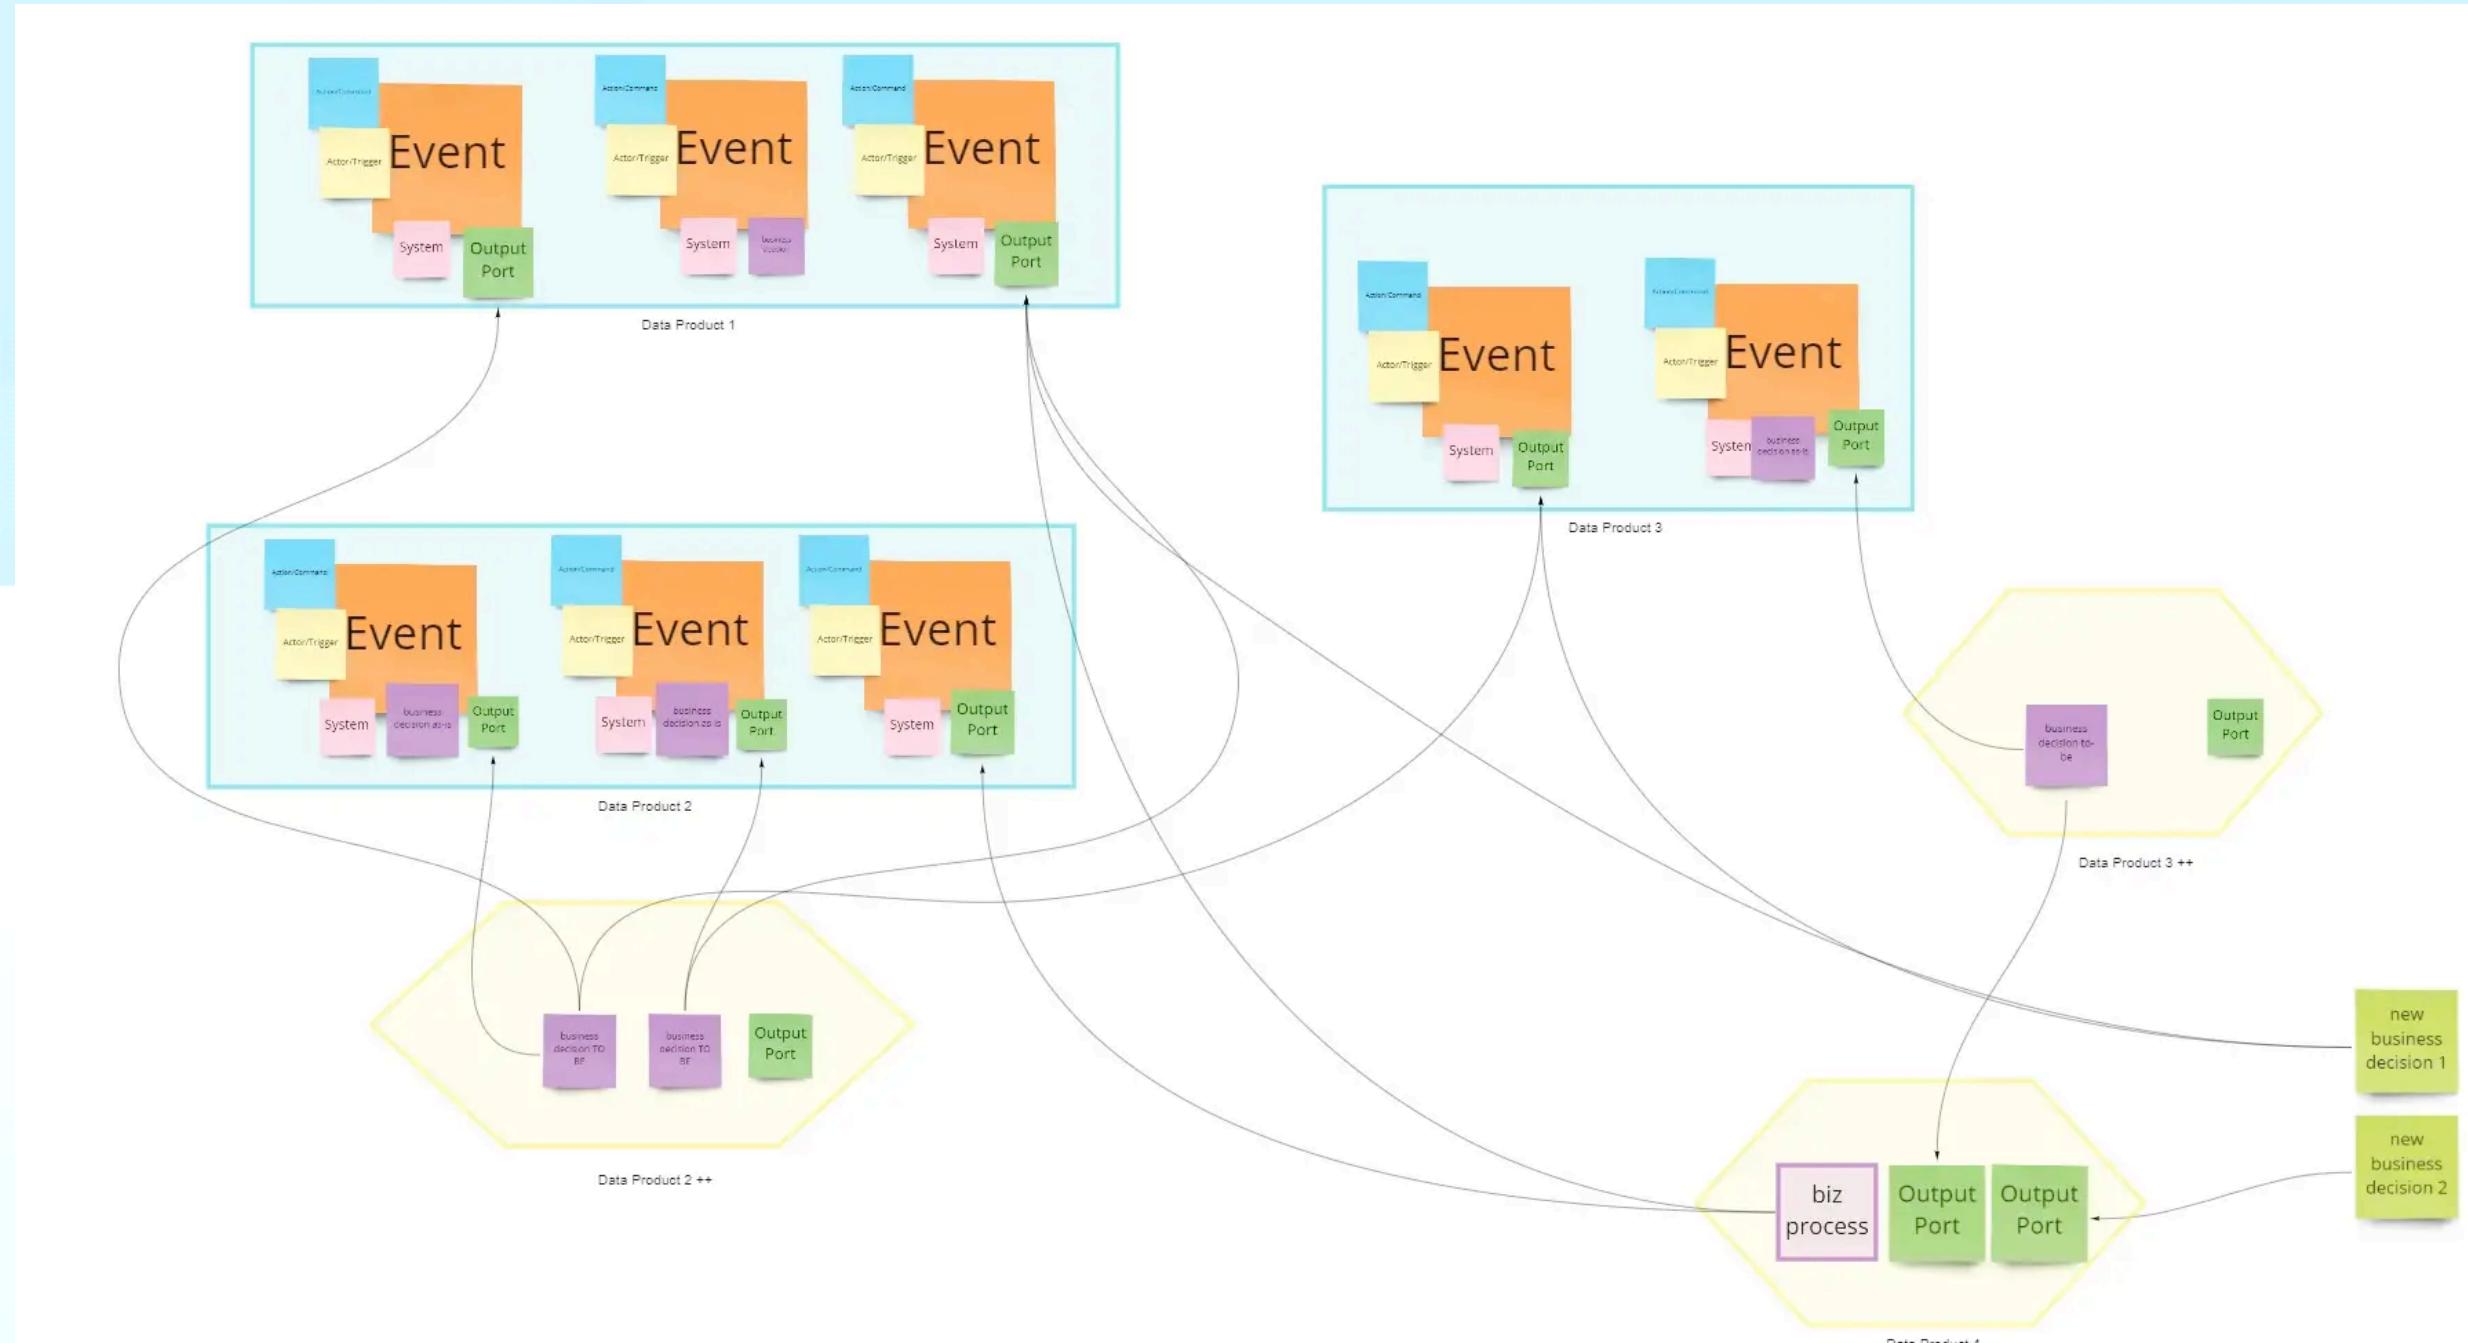
👍 A Spreading Pattern 🔨



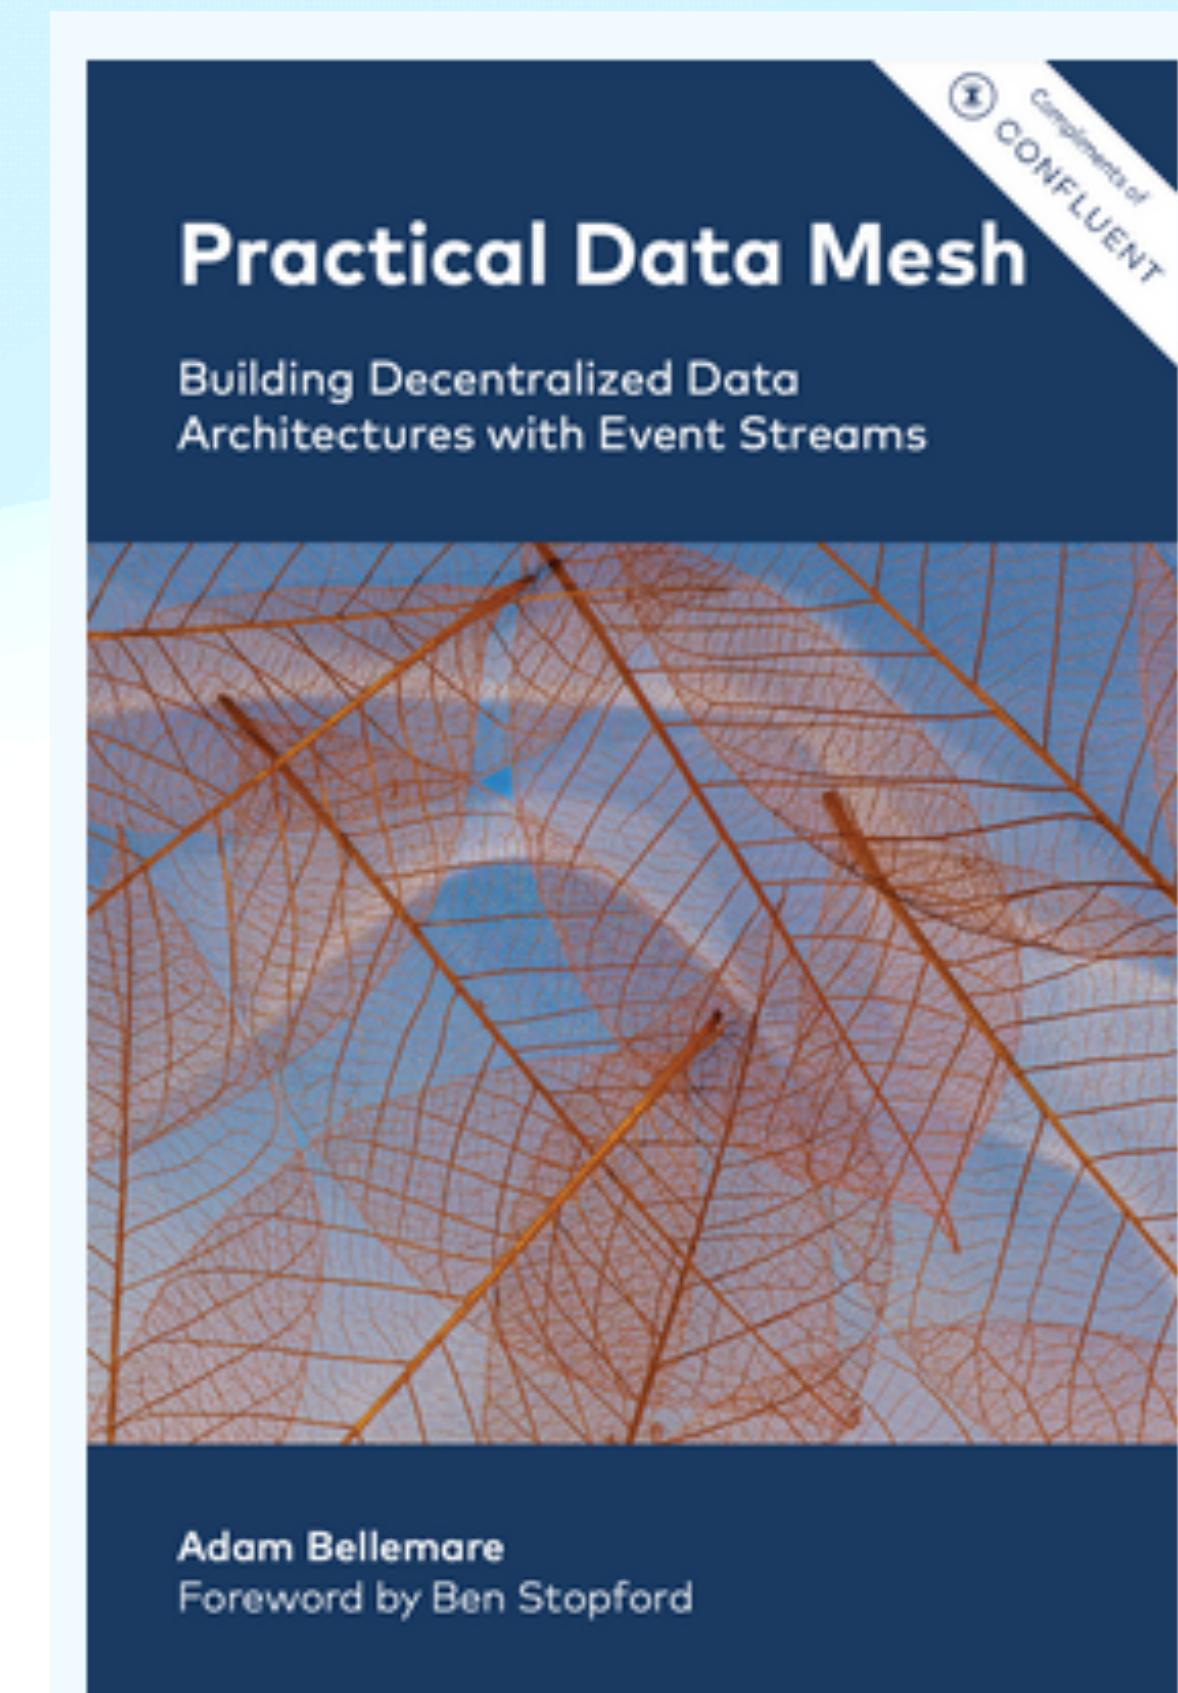
Source: <https://developer.confluent.io/learn-kafka/data-mesh/>

5. Lessons Learned

(Streaming) Data Products



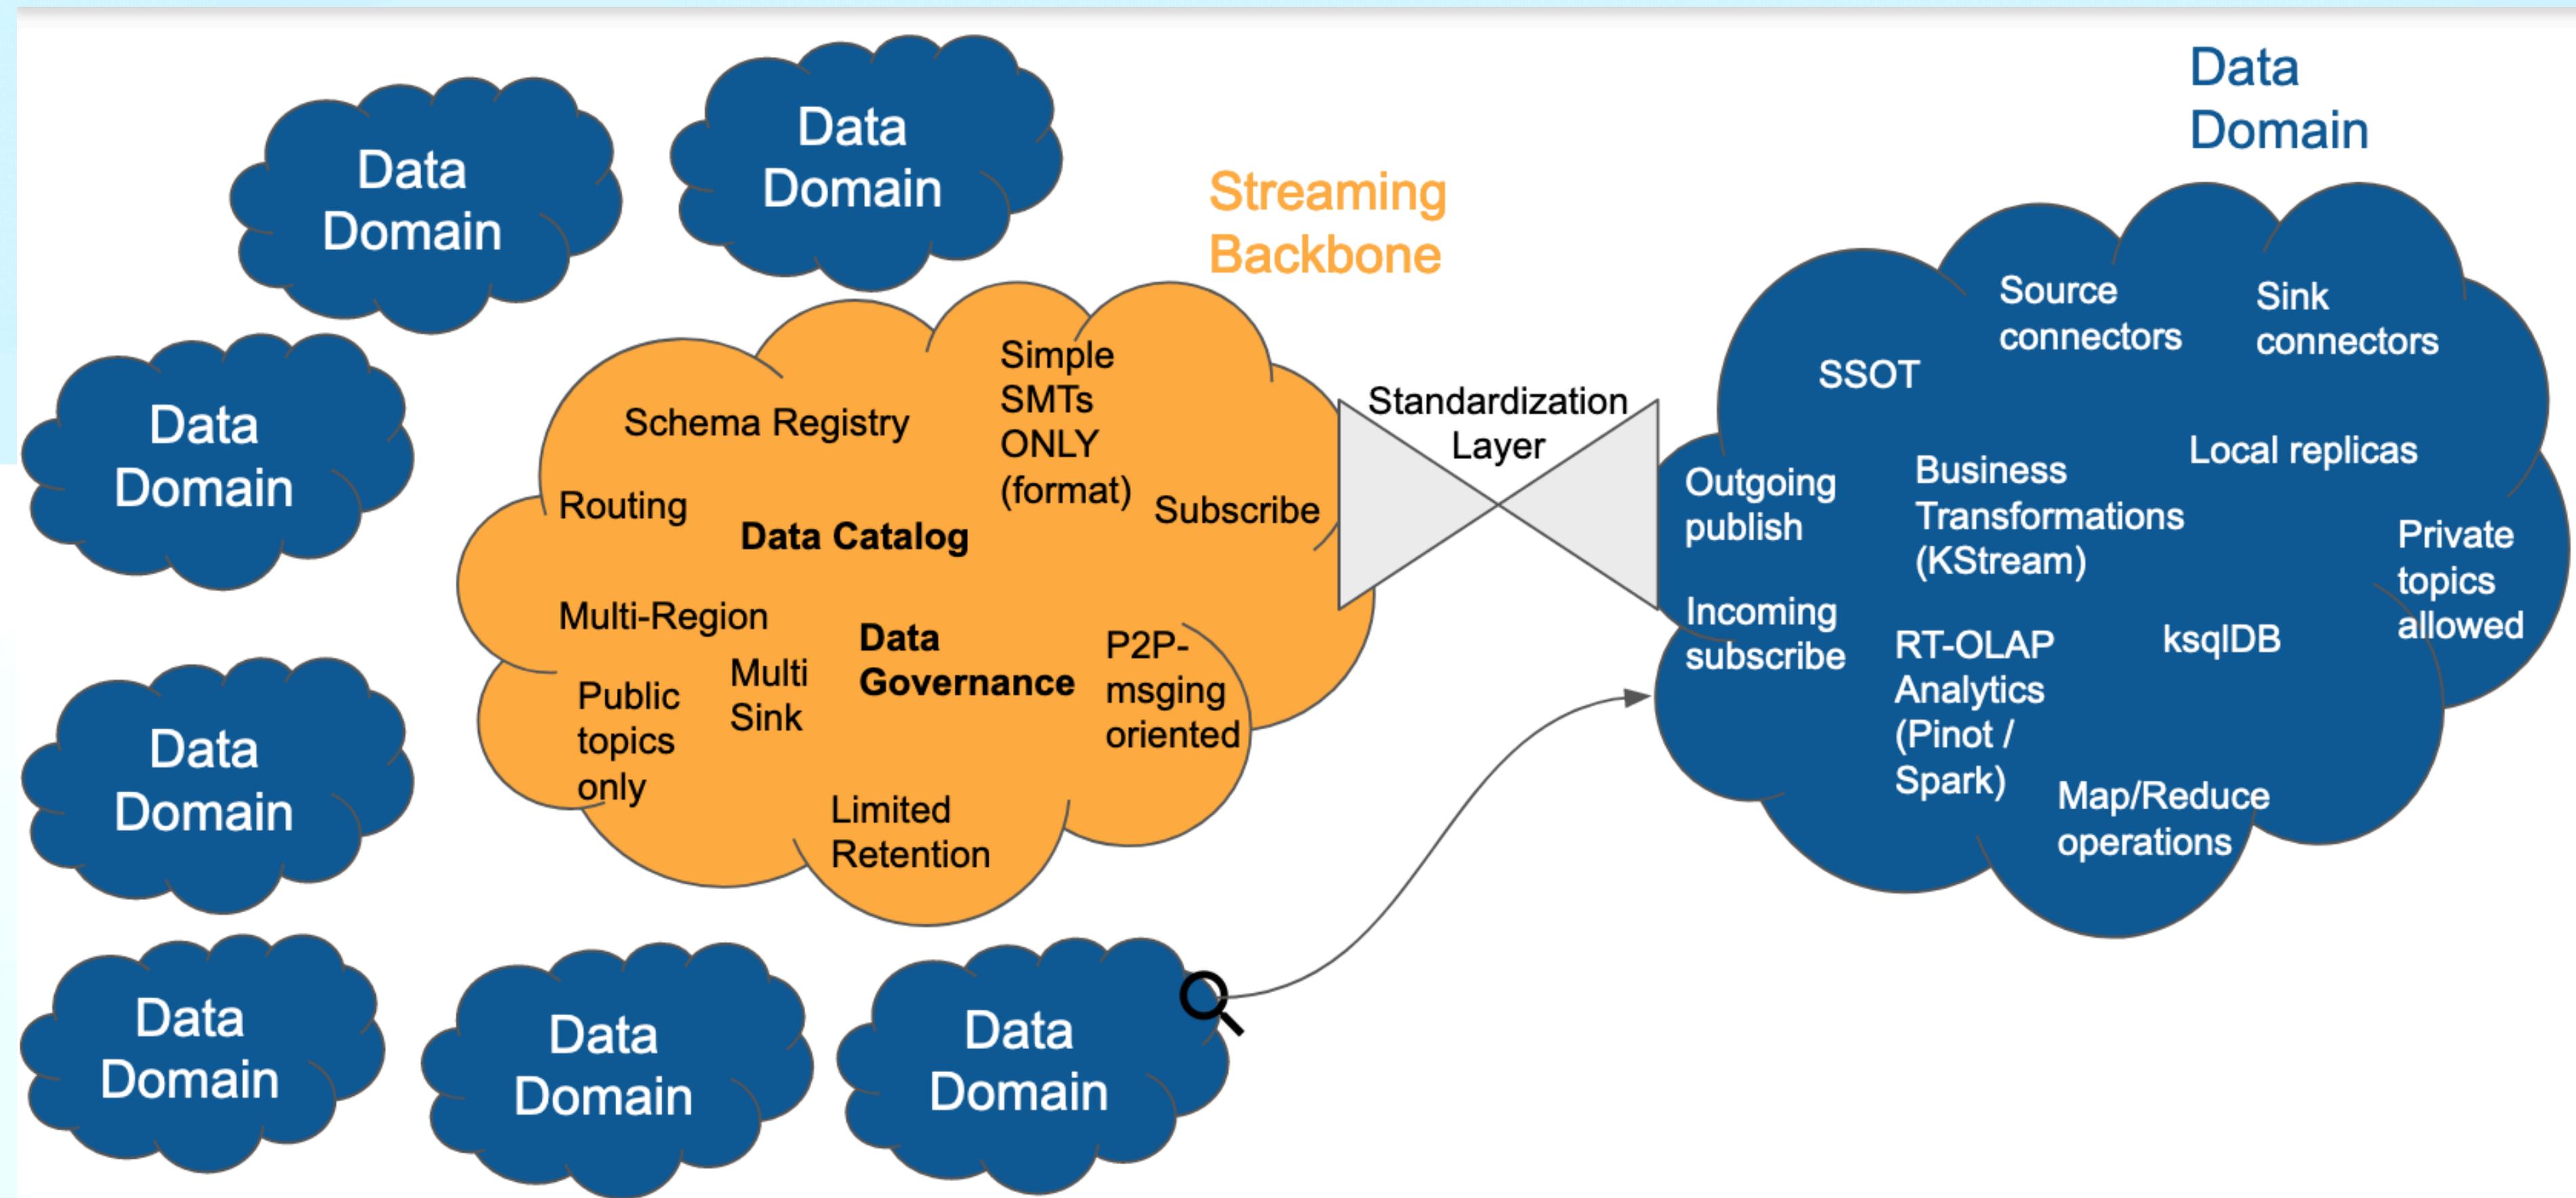
<https://medium.com/agile-lab-engineering/how-to-identify-data-products-welcome-data-product-flow-76d7d85d23af>



<https://developer.confluent.io/learn/data-mesh/>

5. Lessons Learned

👍 A Spreading Pattern 🔨

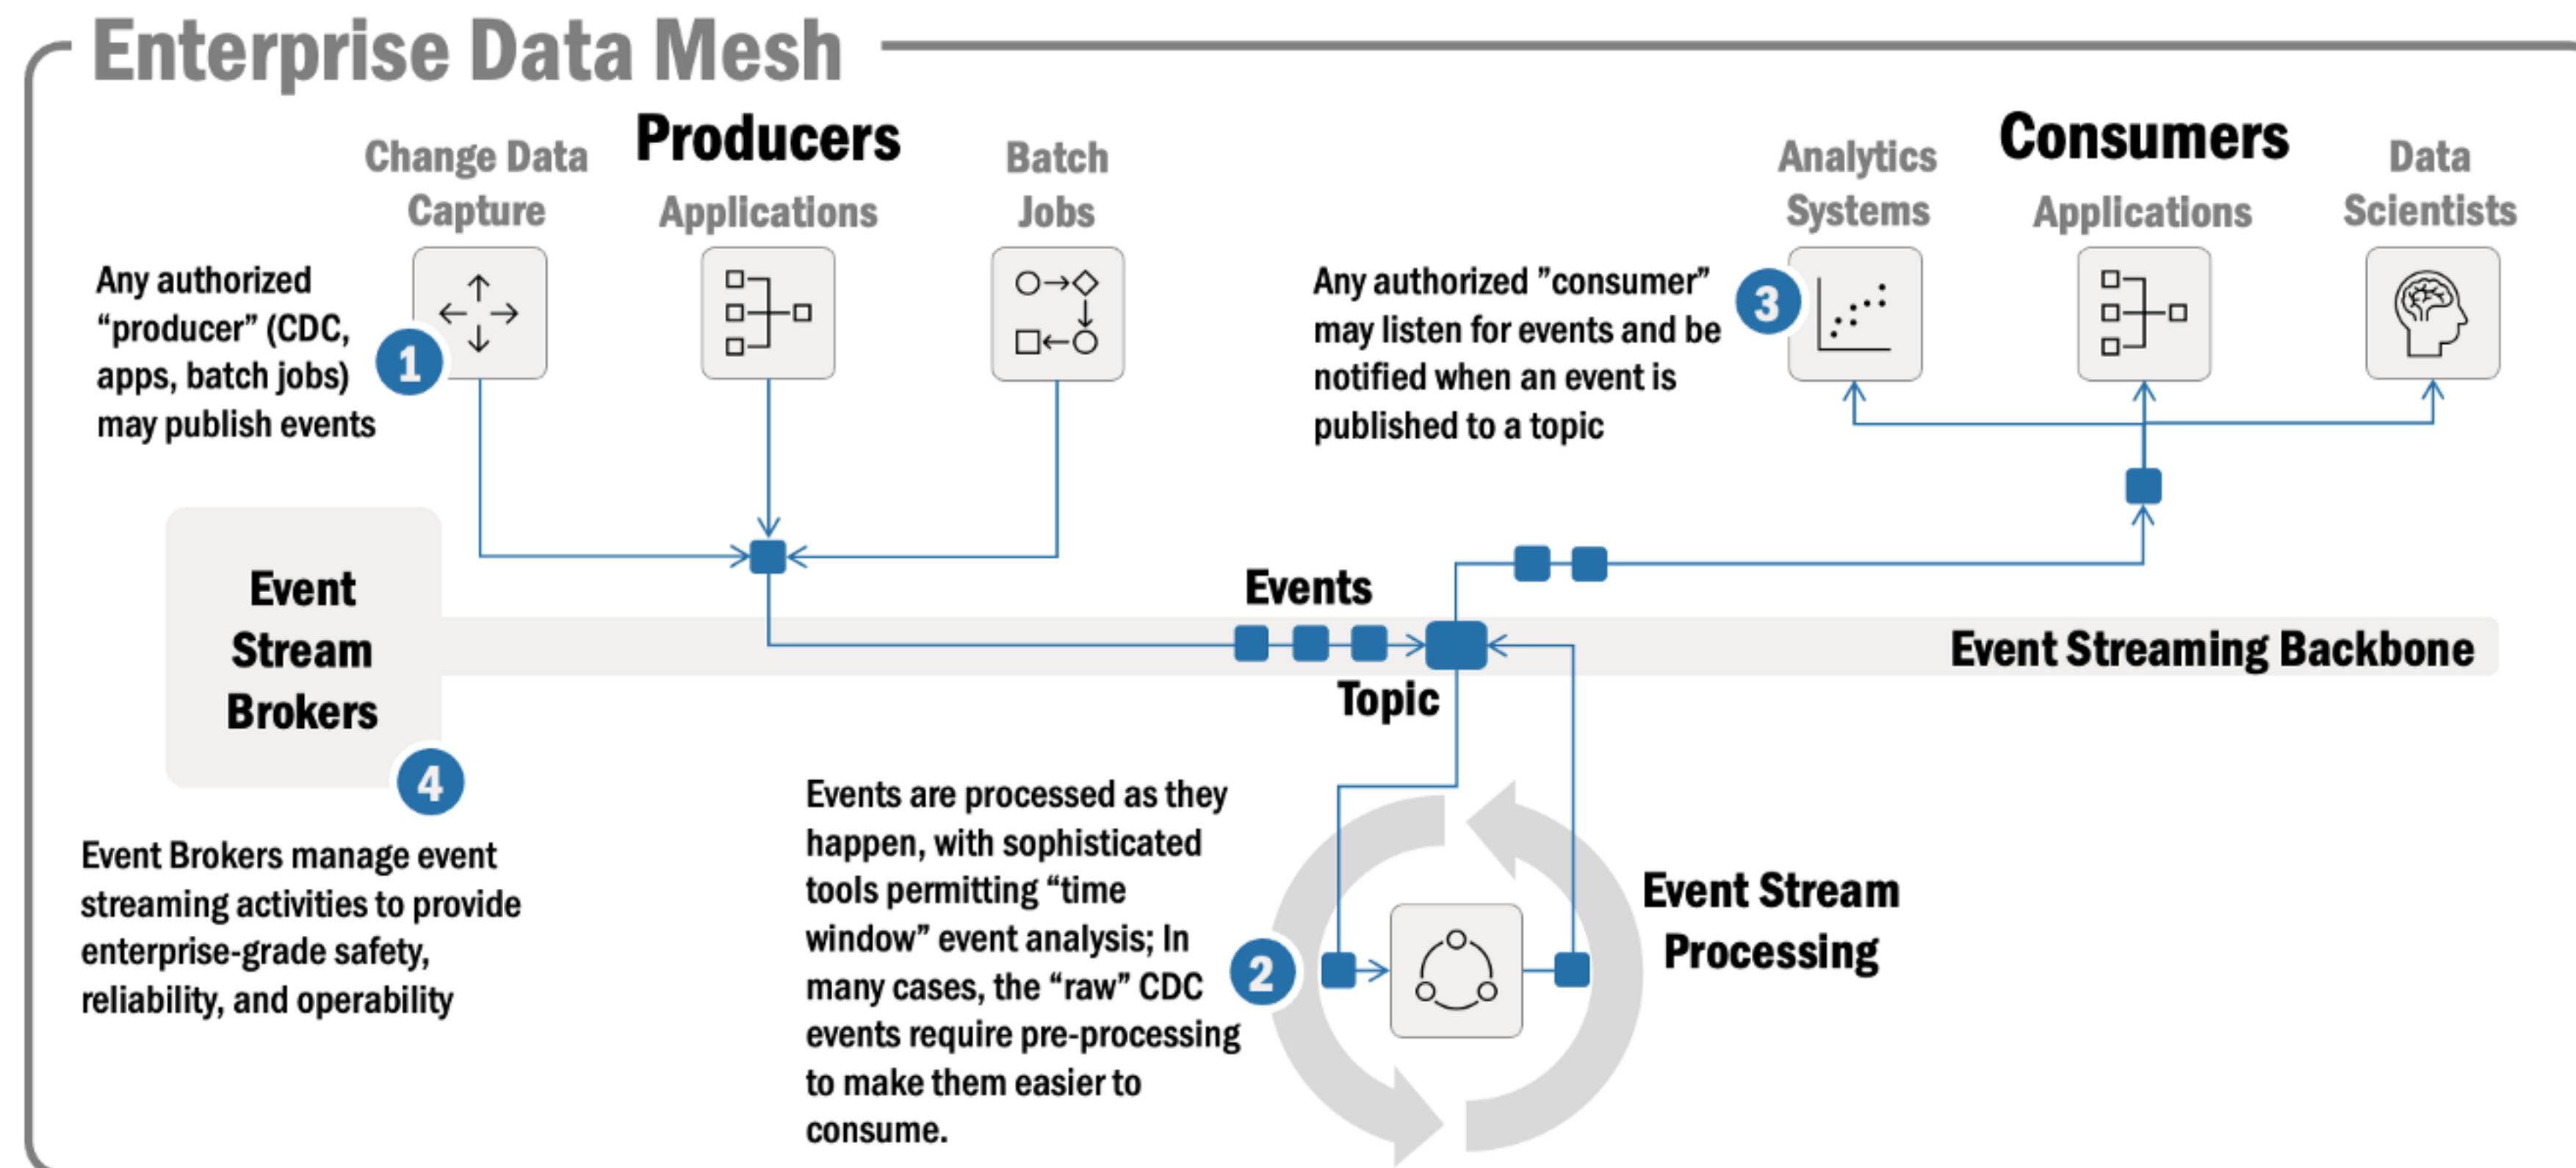


5. Lessons Learned

A Spreading Pattern - (Streaming) Data Mesh

Data Mesh Pattern: Event Streaming Backbone

The “Event Streaming Backbone” pattern is used to transmit data – in near real-time – across the Enterprise Data Mesh



Where do we start from?

Move data,
In the simplest way,
Continuously, better,
Bringing value along the way,
For us first,
=> A Domain

What do we start with?

Events,
Shaping a **source of truth**

How do we do that?

Using Apache Kafka
AND ITS ECOSYSTEM

Where does it lead us?

=> Stream Processing
=> Governance
=> Data (Streaming) **Products**
<= Data (Streaming) **Mesh**

Thank you!

arnaud.esteve@gmail.com

Also: "Apache Kafka: Tips & Tricks I wish I'd known before"